

Building a Design-Based Matching Pipeline: From Principles to Practical Implementation in R

Thomas Leavitt

thomas.leavitt@baruch.cuny.edu

*Marx School of Public and International Affairs
Baruch College, City University of New York (CUNY)
New York, NY 10010, USA*

Luke W. Miratrix

luke__miratrix@gse.harvard.edu

*Harvard Graduate School Of Education
Harvard University
Cambridge, MA 02138, USA*

Abstract

Matching, a canonical design for observational studies, takes many forms that rest on distinct — yet often implicit — statistical principles. We construct a matching pipeline for practitioners that makes these principles explicit and integrates stages often treated separately within a coherent design-based framework. The pipeline begins from the conceptual ideal of a randomized experiment, traces how observational studies depart from it, and then employs matching to approximate that ideal. The next stage is inference under the as-if randomization assumption that matched sets are equivalent to a collection of randomized experiments within blocks, where each block has a fixed number of treated units equal to the number observed in that set. Under this assumption, we consider inference on all individual effects in the "sharp" framework and on the average effect in the "weak" framework. The final stage is a sensitivity analysis to assess, under either framework, how inferences change under departures from as-if randomization. Each step includes extensively commented R code that equips practitioners to implement both established and newly developed procedures, including several not yet available in existing R packages. We illustrate the full workflow through an application examining the effect of United Nations peacekeeping interventions on the duration of post-conflict peace.

1. Design-Based Foundations of the Matching Pipeline

1.1 The Randomized Experimental Ideal

Imagine a randomized experiment in which a researcher flips a fair coin independently for each individual in the study. Heads means assignment of the individual to control, while tails means assignment to treatment. After assignments to treatment and control, the researcher administers the conditions and then compares outcomes between treated and control groups.

Why is this procedure effective? Randomization is a *fair lottery*: Every individual has the same chance of being assigned to treatment. This means that individuals who would respond more strongly to treatment are no more likely to receive it than those who would respond less strongly, and the same is true for control. Because each individual's assignment is determined

by the same coin toss (with the same probability of landing heads or tails), randomization leaves only two possibilities: (1) the difference in outcomes between treatment and control groups reflects the true causal effect, or (2) chance variation produced a misleading difference. Although misleading differences can occur by chance, randomization is valuable because it enables us to use statistical tools to quantify and limit the chance of such errors. As a result, randomized experiments yield especially credible causal conclusions.

Randomization is useful not only as a procedure, but also as an idea. It helps us understand when statistical tools will (and will not) yield credible conclusions, even when we have not directly randomized. In an observational study, the researcher does not control who receives the treatment and instead observes units after they have been assigned to treatment and control groups. In such settings, the idea of randomization can guide how we design studies so that they yield more credible causal conclusions.

1.2 Bridging Randomized Experiments and Observational Studies

A useful framework for connecting randomized experiments to observational studies is what Rubin (1977) calls “assignment to treatment group on the basis of a covariate,” where a covariate is a pre-treatment characteristic of a study’s units. To make this idea concrete, suppose treatment is assigned by independent coin tosses for each individual in which the probability of heads or tails depends on that individual’s value of a single covariate. Consequently, all individuals with the same covariate value share the same chance of ending up in treatment, while those with a different covariate value share a different chance.

In this setting, we can envision forming groups (i.e., matched sets) so that all individuals within a group share the same covariate value. Each group then functions as a miniature randomized experiment in that random chance alone explains why some individuals in the group ended up in treatment while others did not. Importantly, to justify this interpretation, we do not need to know each individual’s actual probability of treatment. It is enough to know that the probability of treatment depends only on the covariate, which ensures that all individuals with the same covariate value have the same chance of treatment.

1.3 Matching to Approximate the Randomized Experimental Ideal

The same intuition applies to matching when the probability of treatment depends on many baseline covariates. The underlying idea is that individuals with similar covariates have similar treatment assignment probabilities. Thus, in an effort to recreate a randomized experiment, we use the covariates we observe and believe determine treatment chances in order to divide individuals into groups, with each group containing both treated and control subjects who are homogeneous in those covariates. We are, in effect, constructing a new single variable: group (i.e., matched set) membership. We can think of this membership variable as the single covariate in the framework of “assignment to treatment group on the basis of a covariate” (Rubin, 1977) discussed above. Although we still do not know each individual’s treatment probability, the hope is that all individuals within a group share the same probability, whatever it may be.

We call the condition in which all individuals within matched sets share the same treatment probability *as-if randomization*, although other terms such as *selection on observables* and *strong ignorability* are common. The term as-if randomization is especially fitting: When all units within each matched set share the same individual treatment probability, all assignments within each set — holding fixed the set’s observed number of treated units — are equally likely. In this case, the probability distribution over these assignments is equivalent to that induced by a randomized experiment in which a researcher first forms blocks corresponding to the matched sets and then randomly assigns a fixed number of units to treatment within each block. Thus, when as-if randomization holds, we can apply the same statistical tools that we would use under such a block-randomized design to draw credible causal inferences from the observational study.

1.4 Why Sensitivity Analysis Matters

Unlike a randomized experiment, even the best matched designs rely on the strong assumption of as-if randomization. When this assumption holds, we can draw causal conclusions by analyzing the data as if they came from a randomized experiment. However, if the assumption is wrong, our causal claims are no longer guaranteed to be credible.

How can this assumption fail? First, individuals within a group may be similar on observed covariates, but not similar enough to have the same treatment probabilities. Second, treatment assignment may depend on covariates we did not measure. If so, even if individuals within groups appear comparable on observed covariates, those individuals may still differ on hidden covariates that determine the probability of treatment.

For these reasons, it is important to assess the sensitivity of our causal conclusions to departures from as-if randomization. Conclusions are especially convincing when they hold not only under this assumption, but also under moderate violations of it. Conclusions that collapse under only mild departures are much less convincing.

1.5 Roadmap of the Design-Based Matching Pipeline

Building on these design-based foundations, we now outline a pipeline that starts with matching and then proceeds to inference and sensitivity analysis:

- **Construct and evaluate matched sets.** Given the specified covariates to be balanced, we begin with the mechanics of optimal matching (Hansen, 2004; Hansen and Klopfer, 2006): choosing a distance measure that defines similarity on the covariates, setting calipers — maximum allowable distances between treated and control units for inclusion in the same matched set — and imposing structural constraints (e.g., requiring matches to be pairs). We then show how to evaluate the resulting design in terms of both effective sample size and covariate balance. For the latter, we focus on tests proposed by Hansen and Bowers (2008) that compare the matched design’s covariate balance to what one would expect under an equivalent completely randomized experiment within blocks (that is, random assignment with a fixed number of treated units per block).

- **Draw causal inferences.** Once a matched design is chosen, practitioners can conduct inference under as-if randomization, either within a “sharp” causal framework, which pertains to individual-level effects for all units, or a “weak” framework, which pertains to a summary quantity of the unit-level causal effects, typically the average treatment effect (ATE). For the sharp framework, we focus on how researchers can use both simulation- and Normal-based approximations to the exact randomization distribution to perform hypothesis tests, which can be inverted to obtain confidence sets and point estimates. For the weak framework, we cover estimation of the ATE and hypothesis tests about it. We emphasize exposition and code for recent variance estimators tailored to designs with only 1 treated or 1 control unit per matched set (Fogarty, 2018; Pashley and Miratrix, 2021) — an important case, since such designs are optimal in terms of balance and effective sample size (Gu and Rosenbaum, 1993; Rosenbaum, 1991; Hansen, 2004).
- **Assess sensitivity.** Finally, we turn to sensitivity analyses under both inferential frameworks. We review established methods for conducting sensitivity analysis for tests of sharp nulls under possible violations of as-if randomization (Rosenbaum and Krieger, 1990; Gastwirth et al., 2000; Rosenbaum, 2018). We then describe and implement new methods that extend sensitivity analysis to tests of weak null hypotheses (Fogarty, 2023).

Below we present a flow diagram summarizing the overall pipeline, partitioned into three sections aligned with the manuscript’s sections. The diagram displays each step and decision point, the relevant R tools (whether existing packages or custom functions included herein), and core references. All relevant code and replication materials are publicly available in the manuscript’s companion GitHub repository (repo).

BUILDING A DESIGN-BASED MATCHING PIPELINE




Figure 1: Flow diagram of the design-based matching pipeline.

2. Implementation of the Matching Pipeline

We provide a high-level overview of the ideas behind matching and include code that demonstrates how to implement those ideas with various R packages. We also work through some of these steps “by hand” to underscore the underlying conceptual issues. Doing so also gives practitioners more flexibility to adapt the matching pipeline to their own needs.

We break the implementation into specific decision points that practitioners commonly face, and present the pipeline in three main stages (see Figure 1):

First Stage: Making a Matched Dataset (toward Justifying As-If Randomization)

- (a) How Do I Measure Similarity on My Chosen Covariates?
 - Similarity on the Estimated Propensity Score
- (b) How Can I Apply Rules for Matches to Ensure Comparability?
- (c) How Can I Apply Rules for Matches to Improve Effective Sample Size?
- (d) How Do I Actually Form the Matches?
- (e) How Do I Decide Whether to Move Ahead with My Matched Design?

Second Stage: Causal Inference (under As-If Randomization)

- (a) How Do I Draw Inferences under the Sharp Framework?
- (b) How Do I Draw Inferences under the Weak Framework?

Third Stage: Sensitivity Analysis (Departures from As-If Randomization)

- (a) How Do Inferences under the Sharp Framework Change under these Departures?
 - Finding the Worst-Case Scenario of Hidden Confounding for Valid Inference
 - Separable Approximation
 - Taylor Series Approximation
 - Conducting Sensitivity Analysis under the Worst-Case Scenario
- (b) How Do Inferences under the Weak Framework Change under these Departures?

Before turning to the first stage, we introduce a running example taken from Gilligan and Sergenti (2008), which we use throughout this manuscript.

2.1 Running Example: United Nations Peacekeeping and Post-Conflict Peace

We introduce matching through an example that examines the causal effect of United Nations (UN) peacekeeping missions on the durability of post-conflict peace, a question of central importance for both academic research and policy. Our example draws on data from Gilligan and Sergenti (2008), whose title includes the phrase “Matching to Improve Causal Inference,” underscoring the value of applying matching to study the UN’s causal impact on post-conflict

peace. These data are publicly available in the supplementary information of the article’s webpage in the *Quarterly Journal of Political Science* (DOI: 10.1561/100.00007051).

The dataset we use from Gilligan and Sergenti (2008) includes 87 observations, each corresponding to a country’s peace period episode following a civil war, with episodes beginning as early as January 1989 and data extending through December 2003. In some episodes, UN peacekeepers intervened (e.g., Sierra Leone, Jan 2001 - Dec 2003), while in others they did not (e.g., Macedonia, Sep 2001 - Dec 2003). The treatment variable is UN intervention (UN), coded as 1 if a UN mission was present during the peace period and 0 otherwise. The outcome variable is the duration of the peace spell, which Gilligan and Sergenti (2008) measure as the log of the number of days from the start of peace until either the outbreak of a new conflict or right-censoring at December 2003. This log-transformed outcome (`ldur`) captures the total length of the peace period rather than the time elapsed after a potential UN intervention. In practice, however, UN interventions almost always began immediately after the onset of peace: “Of the 19 post-conflict UN interventions, the United Nations was present within the first month for 16 of them” (Gilligan and Sergenti, 2008, p. 118). In what follows, we set aside these measurement details.

As Gilligan and Sergenti (2008) state, “UN missions are not randomly assigned” (p. 89). Whether peacekeepers are present during a given post-conflict peace period presumably depends on baseline covariates measured by Gilligan and Sergenti (2008), including the logged number of deaths (`lwdeaths`), the logged duration of the previous war (`lwdurat`), ethnic fractionalization (`ethfrac`, a 0-100 index for which dividing by 100 is intended to represent the probability that two randomly chosen individuals belong to different ethnic groups), logged population size (`pop`) and others. We implement matching using these same covariates, but emphasize that our exercise is expository and not intended as a replication of the original findings.

2.1.1 LOADING THE DATA FOR THE RUNNING EXAMPLE

To load the data, you could first download the replication files from the supplementary information on the article’s webpage (DOI: 10.1561/100.00007051), save the files to your working directory, and then use a package such as `haven` to load the Stata (.dta) file, `peace_pre_match.dta`, into R. However, for our purposes, we recommend loading our pre-created .rds file (`peace_pre_match.rds`), in which Stata’s monthly numeric dates have been converted to R’s year-month format and the geographic region indicators recoded into a single factor variable (`region`). The command below loads this pre-created dataset (`peace_pre_match.rds`) into the R environment as an object named `data`.

```

1 # Define base URL for the Matching Guide GitHub repo
2 base_url <- "https://raw.githubusercontent.com/tl2624/matching-guide/main"
3
4 # Load the cleaned dataset
5 data <- readRDS(
6   url(
7     paste0(
8       base_url, "/data/peace_pre_match.rds"
9     )
10  )

```

11)

2.2 First Stage: Making a Matched Dataset (toward Justifying As-If Randomization)

In an observational setting, treatment is not assigned by the flip of a coin but depends on individuals' covariates. The goal of matching is to compare treated and control units that have the same chances of treatment based on those covariates — i.e., the same *propensity scores*. If we could observe propensity scores, it would be straightforward to compare treated and control units by matching them directly on their propensity scores. Because propensity scores are not directly observed, we instead aim to create a collection of matched sets that is *balanced* — meaning that treated and control observations are similar in their covariates.

Before turning to questions of covariate similarity and matching, it is important to note that substantive transformations of covariates play a central role, as they determine the inputs on which subsequent notions of similarity between units are based. These transformations reflect substantive judgments and are often among the most important decisions in practice. In this manuscript, however, we do not address these substantive choices. Instead, we take the transformations used by Gilligan and Sergenti (2008) as given. For example, they measure ethnic fractionalization (`ethfrac`) on a 0-100 scale (rather than a 0-1 scale) and apply logarithmic transformations to covariates such as the number of deaths (`lwdeaths`) and the duration of the previous war (`lwdurat`), among others. Our focus is on design and analysis decisions conditional on these substantively chosen scales.

2.2.1 HOW DO I MEASURE SIMILARITY ON MY CHOSEN COVARIATES?

In the simplest terms, matching is about ensuring apples-to-apples, rather than apples-to-oranges, comparisons between treated and control observations (Rubin and Waterman, 2006). To create matched sets in which treated and control groups are similar in their covariates, we first need a distance measure that quantifies how close any two observations are. With such a measure in hand, we can then construct sets of treated and control observations that are close on this measure — i.e., apples-to-apples in their pre-treatment characteristics.

We record the distances between treated and control units in a distance matrix: The rows correspond to treated units and the columns to control units. Each cell of the matrix records the distance between a specific treated unit and a specific control unit, as defined by a distance measure. This distance measure takes the baseline covariates of the two units and maps them to a single nonnegative number, with smaller values indicating greater similarity.

In our setting, we are interested in similarity across 9 covariates, named in the object `covs`.

```
12 # Define character vector of the 9 covariate names in the dataset
13 covs <- c("lwdeaths", "lwdurat", "ethfrac", "pop", "lmtnest", "milper",
14           "bwgdp", "bwplty2", "region")
```

The first four — `lwdeaths`, `lwdurat`, `ethfrac`, and `pop` — were introduced earlier. The others include a logged measure of the proportion of a country's land area that is mountainous (`lmtnest`), the logged total number of military personnel in a country (`milper`), logged GDP per capita before the last civil war (`bwgdp`), the Polity score (a -10 to 10 scale, from

autocratic to democratic institutions) before the last civil war (`bwplty2`), and a region factor (`region`) with categories for Eastern Europe, Latin America, Asia, Sub-Saharan Africa, and North Africa/Middle East.

All of these covariates are measured before treatment and presumably determine the chance of a UN intervention during a country's peace period. Our interest in them stems primarily from their role in determining those intervention probabilities. Yet many of these covariates may also be prognostic; that is, they help predict the outcome of interest in Gilligan and Sergenti (2008), the log duration of the peace period (`ldur`) that countries would potentially experience with or without a UN intervention. This prognostic value of covariates can provide an additional reason to match on them (Hansen, 2008a; Sales et al., 2018).

There are many ways to measure the distance between a treated and a control unit. For example, we might compare units using the Euclidean distance — i.e., the square root of the sum of squared differences — across all baseline covariates. To do so, we first convert the factor variable `region`, which stores categorical labels, into a set of dummy (0 or 1) variables for each region. This conversion ensures that distances between treated and control units can be computed, since distance measures require numeric variables rather than categorical labels.

```

15 # Install "dplyr" package (only run if not already installed)
16 # install.packages("dplyr")
17
18 # Load package for data manipulation (mutate, group_by, summarize, etc.)
19 library(dplyr)
20
21 # Convert categorical variable "region" into 0/1 dummy indicators
22 data <- data |> # Pipe (|>) to pass left-hand result into next function call
23   mutate(
24     # Use mutate() to create new variables by transforming existing columns
25     # ifelse() returns one value if condition is TRUE and another if FALSE
26     eeurop = ifelse(test = region == "eeurop", yes = 1, no = 0),
27     lamerica = ifelse(test = region == "lamerica", yes = 1, no = 0),
28     asia = ifelse(test = region == "asia", yes = 1, no = 0),
29     ssafrica = ifelse(test = region == "ssafrica", yes = 1, no = 0),
30     nafrme = ifelse(test = region == "nafrme", yes = 1, no = 0)
31   )
32
33 # Covariate names with "region" replaced by dummy indicators
34 expanded_covs <- c(
35   # setdiff() returns elements in 'x' that are not in 'y'
36   setdiff(
37     x = covs,
38     y = "region" # Drops "region" from the covariate list
39   ),
40   "eeurop", "lamerica", "asia", "ssafrica", "nafrme"
41 )

```

As an example, we calculate the Euclidean distance between post-conflict Liberia, where the UN did intervene, and post-conflict Guinea-Bissau, where the UN did not.

```

42 # Extract covariates for Liberia and Guinea-Bissau (cname = country name)
43 liberia <- data[data$cname == "Liberia", expanded_covs]

```

```

44
45 guinea_bissau <- data[data$cname == "Guinea-Bissau", expanded_covs]
46
47 # Compute Euclidean distance between the two countries on these covariates
48 sqrt(
49   sum(
50     (liberia - guinea_bissau)^2
51   )
52 ) # Display the Euclidean distance

```

```
[1] 57.44263
```

We can obtain the full matrix of pairwise distance values using `match_on()` from the `optmatch` package. We do not need to manually recode factor variables into dummy indicators because `match_on()` handles this conversion automatically.

```

53 # Create a formula: UN (treatment indicator) ~ covariates
54 # Note: we keep "region" in covs as a factor
55 cov_fm1a <- reformulate(
56   termlabels = covs,
57   response = "UN"
58 )
59
60 # Install optmatch if not already installed
61 # install.packages("optmatch")
62
63 # Load optmatch, which provides the match_on() function
64 library(optmatch)
65
66 # Euclidean distance matrix: treated (UN = 1) vs control (UN = 0)
67 dist_mat_euc <- match_on(
68   x = cov_fm1a,           # Formula for covariates
69   data = data,           # Dataset used
70   standardization.scale = NULL, # No rescaling of covariates
71   method = "euclidean"   # Use Euclidean distance
72 )
73
74 # Add country names (cname) as row/column labels for clarity
75 dimnames(dist_mat_euc) <- list(
76   data$cname[data$UN == 1], data$cname[data$UN == 0]
77 )
78
79 # Display submatrix of distances:
80 round(
81   # Round distances for display
82   # treated units in rows 11-15 vs. control units in columns 27-30
83   x = dist_mat_euc[11:15, 27:30],
84   digits = 2 # Number of decimal places to round
85 )

```

	Niger	Guinea	Togo	Central African Republic
Sierra Leone	94.21	101.18	116.32	116.94
Zaire	26.68	29.87	45.63	46.57
Rwanda	66.75	71.57	77.04	76.03
Mozambique	133.47	140.60	155.40	155.80

Namibia 246.44 253.23 268.35 268.08

In this subset of the full distance matrix, the first entry is the Euclidean distance of Sierra Leone (treated) from Niger (control). The last listed entry is the distance between Namibia (treated) and the Central African Republic (control).

One concern with Euclidean distance is that it depends on the scale of the covariates. For example, the difference between a country in Sub-Saharan Africa (`ssafrica` = 1) and a country in Latin America (`ssafrica` = 0) would contribute the same to the Euclidean distance as the difference between two countries with GDP per capita values of \$3000 and \$3001. Intuitively, we would not want such a tiny difference in economic size to be treated as equally important as belonging to different regions of the world. More generally, we want differences across variables to be placed on a comparable scale, so that a meaningful difference in one variable counts about the same as a difference of similar importance in another.

Another concern with Euclidean distance is that it ignores correlations among covariates. For example, countries with larger populations (`pop`) usually have more military personnel (`milper`), if only because a larger population provides a greater pool of potential recruits. Therefore, differences in both covariates (`pop` and `milper`) may largely reflect the same underlying factor — population size (`pop`). Yet Euclidean distance adds these differences separately, as if they were unrelated, which can exaggerate the overall distance between two observations.

The Mahalanobis distance (Mahalanobis, 1936) addresses both of these concerns. First, it standardizes covariates so that differences are placed on a comparable scale. Second, the Mahalanobis distance adjusts for correlations among covariates, ensuring that highly related variables are not effectively counted twice. We can construct a distance matrix based on Mahalanobis rather than Euclidean distances as follows.

```
85 # Mahalanobis distance matrix: treated (UN = 1) vs control (UN = 0)
86 dist_mat_mah <- match_on(
87   x = cov_fmla,
88   data = data,
89   standardization.scale = NULL,
90   method = "mahalanobis" # Use Mahalanobis distance
91 )
```

The standardization in the Mahalanobis distance is distinct from researchers' substantive choices about how to transform covariates. Substantive transformations — such as logging certain covariates — reflect subject-matter considerations and precede the matching procedure. The standardization in the Mahalanobis distance is a statistical device that operates on the pre-specified covariates to compute distances comparably across covariates. The purpose of this device is to facilitate matches that are similar on the original covariate scales.

Similarity on the Estimated Propensity Score So far, we have focused on ways to measure distance between treated and control units across many covariates in order to identify which units are most similar and group them together to achieve covariate balance. However, when there are many covariates, it becomes difficult to find treated and control

units that are similar on all covariates. This challenge is often referred to as the “curse of dimensionality.”

A common way to address this problem is to reduce the information from many covariates into a lower-dimensional form. The *estimated propensity score* does this by collapsing information from all covariates into a single number. This number represents a transformation of an estimated linear index of covariates that accounts for how strongly each covariate predicts treatment.

Consider, for example, a logistic model for the estimated propensity score of an individual unit i . We write this model as

$$\hat{\lambda}(\mathbf{x}_i) = \frac{1}{1 + \exp(-\hat{\boldsymbol{\beta}}^\top \mathbf{x}_i)}, \quad (1)$$

where $\mathbf{x}_i$ is the covariate vector, $\hat{\boldsymbol{\beta}}$ is the vector of estimated coefficients, and the superscript $^\top$ denotes the transpose, turning the row vector $\hat{\boldsymbol{\beta}}$ into a column vector. The quantity $\hat{\boldsymbol{\beta}}^\top \mathbf{x}_i$ is the estimated linear index, and the inverse logistic function, $1/(1 + \exp(-x))$, maps any real-valued input, x , onto the interval $(0, 1)$. The estimated linear covariate index for unit i , $\hat{\boldsymbol{\beta}}^\top \mathbf{x}_i$, is simply the logit, i.e., log odds, transformation of $\hat{\lambda}(\mathbf{x}_i)$ in (1).

This quantity, $\hat{\lambda}(\mathbf{x}_i)$, is a simple transformation of a linear index of covariates that best “separates” treated from control units. The estimated coefficients ($\hat{\boldsymbol{\beta}}$) “separate” treated from control units on the linear index because the coefficients are chosen to maximize a likelihood that rewards large differences between the groups. When treated and control units have little or no covariate overlap, the estimated linear indices can diverge substantially, very positive for treated units and very negative for controls. With greater overlap, the estimated linear indices for treated and control units are similar, clustering near zero.

This estimation process reflects how predictive each covariate is of treatment. When treated and control observations lack overlap on a covariate, that covariate is highly predictive of treatment and therefore receives an estimated coefficient with a large magnitude. When there is substantial overlap on a covariate, it is less predictive of treatment, and the magnitude of its coefficient is small. Consequently, when we assess similarity on the estimated linear index of covariates in (1), the estimated coefficients assign greater importance to covariates that strongly predict treatment and less importance to those that do not.

To see this logic in action, we first estimate a logistic propensity score model using all covariates except `region`, which we exclude because some regions perfectly — or nearly perfectly — predict treatment assignment, leading to (near-)complete separation (Albert and Anderson, 1984).

```

92 # Formula for UN ~ covariates (excluding "region")
93 psm_cov_fm1a <- reformulate(
94   termlabels = setdiff(
95     x = covs,
96     y = "region"
97   ),
98   response = "UN"
99 )
100
```

```

101 # Fit logistic regression for propensity score model
102 psm <- glm(
103   formula = psm_cov_fm1a,           # Treatment ~ covariates
104   family = binomial(link = "logit"), # Logistic regression (logit link)
105   data = data
106 )

```

Using this fitted propensity score model, we can extract each unit’s estimated linear covariate index, $\hat{\beta}^\top x_i$, as follows.

```

107 # Extract logit propensity scores (linear predictors from fitted model)
108 lin_cov_inds <- psm$linear.predictors # = model.matrix(psm) %*% coef(psm)

```

Below we can see that the estimated linear covariate indices from the model `psm` correspond exactly to the logits (i.e., the log-odds transformations) of the model’s predicted probabilities of treatment, which range between 0 and 1.

```

109 # Extract estimated propensity scores (predicted probabilities of UN = 1)
110 p_scores <- psm$fitted.values
111
112 # Check that lin_cov_inds equals log odds (propensity scores on logit scale)
113 all.equal(
114   lin_cov_inds, log(p_scores / (1 - p_scores))
115 )
116
117 # Also check that p_scores equals logistic(lin_cov_inds)
118 all.equal(
119   p_scores, 1 / (1 + exp(-lin_cov_inds))
120 )

```

To see how two observations that differ on many covariates can still have similar estimated propensity scores, consider post-conflict Namibia (treated) and Burundi (control). The two are similar on some covariates, such as logged military personnel (`milper`), but — consistent with the “curse of dimensionality” — very different on others, such as duration of the last war (`lwdurat`) and ethnic fractionalization (`ethfrac`). Yet the estimated linear indices for Namibia and Burundi differ by only a minuscule amount. This small difference occurs because the covariates on which Namibia and Burundi differ greatly have small estimated coefficients in the propensity score model (e.g., approximately -0.01 for `lwdurat` and approximately 0 for `ethfrac`), while those on which the two observations are similar, such as `milper`, have large estimated coefficients (approximately -0.92).

To illustrate these broader patterns, the boxplot below compares the distributions of the estimated linear covariate index for treated and control groups.




Figure 2: The empirical distributions of the linear covariate index for treated and control groups.

As the figure above shows, there is some, but not a lot of, overlap between treated and control groups. In accordance with our earlier discussion of how the estimated linear covariate index “separates” treated from control units, many control observations have very negative values (as low as -7.52), far from the treated units’ range of -2.24 to 2.21. Nevertheless, a sufficient number of treated and control observations have estimated linear covariate indices that cluster around 0, indicating covariate overlap for at least a subset of treated and control observations.

2.2.2 HOW CAN I APPLY RULES FOR MATCHES TO ENSURE COMPARABILITY?

The discussion above on measuring covariate distances between treated and control observations helps identify which treated-control pairs are similar. The goal of identifying these similar treated-control pairs is to form matched sets that are closely aligned on their covariates, thereby improving balance between treatment and control groups. In practice, we do this by excluding potential matches that are “too dissimilar.” There are two common ways to do this:

- **Exact matching:** Require that units be identical on some subset of important covariates, typically those that are categorical or coarse enough for units to take the same values.

- **Calipers:** Impose a threshold for the maximum allowable distance so that no matched set may include a treated and a control unit that are farther apart than this caliper.

As a simple example to build intuition, we will impose the following constraints:

- Observations can belong to the same matched stratum only if they are in the same geographic region.
- Treated and control observations more than two points apart on the Polity score cannot be in the same matched set.

Below we impose the first constraint, requiring an exact match on region. Doing so produces separate distance matrices containing the Euclidean distances on the covariate used to define the exact match (`region`), where all entries are 0, indicating that treated and control units belong to the same region.

```
121 # Create distance structure: 0 if units are in same region, Inf otherwise
122 em_region <- exactMatch(
123   x = UN ~ region,
124   data = data
125 )
```

Now we impose the second constraint: We construct a distance matrix based on the Polity score, `bwplty2`, with a caliper of 2. This matrix records the Euclidean difference in Polity scores when the difference is 2 or less, and assigns a value of ∞ (denoted in R as `Inf`) when the difference exceeds 2. The ∞ entries are crucial because `optmatch` minimizes the sum, across all matched sets, of the within-set sums of covariate distances between each treated-control pair. Consequently, any treated-control pair differing by more than 2 points on the Polity score is assigned an overall distance of ∞ , which prevents them from being matched.

```
126 # Euclidean distance on Polity score (bwplty2) with caliper = 2
127 # Pairs differing by > 2 are set to Inf
128 euc_dist_polity_cal_2 <- match_on(
129   x = UN ~ bwplty2,
130   caliper = 2, # Set caliper
131   data = data,
132   standardization.scale = NULL,
133   method = "euclidean"
134 )
```

Note that we used Euclidean distance here because, after exactly matching on geographic region, matching proceeds on only 1 covariate (Polity score), so we do not have to worry about covariates' relative scales.

Finally, we combine the two constraints into distance matrices defined within each region by “adding” the objects `em_region` and `euc_dist_polity_cal_2`, as shown below.

```
135 # Overall distance matrix by element-wise addition of two distance matrices
136 overall_dist_mat <- em_region + euc_dist_polity_cal_2
```

2.2.3 HOW CAN I APPLY RULES FOR MATCHES TO IMPROVE EFFECTIVE SAMPLE SIZE?

Beyond comparability in covariates, we also care about the matched study’s size. The effective sample size depends not simply on the total number of units included in our matches. Effective sample size also depends on how those units are arranged across the matched sets.

The `optmatch` package will always produce matches with a particular arrangement of units across sets. In particular, all sets contain either 1 treated unit or 1 control unit — an overall structure that minimizes imbalance while excluding as few units as possible (Rosenbaum, 1991; Gu and Rosenbaum, 1993; Hansen, 2004). In practice, optimal full matching can yield lopsided sets, with either 1 treated matched to many controls or 1 control matched to many treated units, which has implications for the effective sample size.

In the `optmatch` package, the effective sample size is defined as the sum, across matched sets, of the harmonic mean of the numbers of treated and control units within each set. The formal definition is given by

$$\textbf{Effective Sample Size: } \sum_{s=1}^S \left[\left(m_s^{-1} + (n_s - m_s)^{-1} \right) / 2 \right]^{-1}, \quad (2)$$

where the index s runs over the $\{1, \dots, S\}$ matched sets, with m_s denoting the number of treated units in set s and $n_s - m_s$ the number of control units. With n_s denoting the number of units in set s , the total number of individuals included in the matched study is $n = \sum_{s=1}^S n_s$.

From the formula in (2), we can see how the effective sample size depends on the arrangement of units across sets. For example, in a study with 4 total units, the effective sample size would be 2 if the units were arranged into 2 matched pairs. By contrast, in a study of the same total size but arranged as a single set with 1 treated unit and 3 controls, the effective sample size would be 1.5. The effective sample size is larger in the former arrangement because it provides two distinct treated-versus-control comparisons, whereas the latter provides only one.

This definition of effective sample size anticipates subsequent outcome analysis under as-if randomization, as effective sample size bears directly on the precision of estimators and the power of hypothesis tests. Assuming constant, additive treatment effects within a set, the variance of the Difference-in-Means — the average outcome among treated units minus the average outcome among control units — in that set is minimized when the harmonic mean is largest (Hansen and Bowers, 2008; Hansen, 2011). The harmonic mean reaches its maximum when the numbers of treated and control units are equal. Thus, all else equal, matched pairs and other balanced sets provide more information about causal effects than lopsided sets with unequal treated-to-control ratios.

One straightforward way to increase effective sample size is to relax restrictions on which units can be matched. Relaxing calipers can admit additional units into the matched design by allowing matches for units that would otherwise be discarded. For instance, we might widen the caliper on Polity score from 2 to 3 and then rebuild the distance matrix.


```

137 # Apply a caliper of width 3 to the polity Euclidean distance matrix
138 euc_dist_polity_cal_3 <- match_on(
139   x = UN ~ bwplty2,
140   caliper = 3,
141   data = data,
142   standardization.scale = NULL,
143   method = "euclidean"
144 )
145
146 # Combine regional exact match distance with polity distance
147 em_region + euc_dist_polity_cal_3

```

Doing so increases the effective sample size, but such gains need not be substantial. When newly admitted units are absorbed into existing matched sets — each of which contains only one treated or one control unit — rather than generating additional treated-versus-control comparisons, the resulting contribution to effective sample size is limited. More generally, as the formula in (2) makes clear, the effective sample size depends on how units are distributed across sets: When units, including newly admitted ones, are concentrated in only a few sets, the resulting gain in effective sample size is modest.

Instead of relying solely on caliper width, researchers can also shape the effective sample size by controlling the structure of matched sets via the `min.controls` and `max.controls` arguments in `optmatch`'s `fullmatch()` function. These arguments specify lower and upper bounds, respectively, on the ratio of control to treated units within each matched set. By default, `min.controls = 0` and `max.controls = Inf`, which impose no restrictions on matched set composition. Changing these defaults allows researchers to control how balanced or lopsided matched sets may be, thereby shaping the effective sample size.

To illustrate, suppose we want to restrict matches using the `overall_dist_mat` introduced earlier. We know that `optmatch` will divide the data into matched sets containing one treated unit and any positive number of controls, or one control unit and any positive number of treated units. However, we can impose additional constraints on this full matching, such as requiring a minimum control-to-treated ratio of 1:2 — that is, at least one control for every two treated units (`min.controls = 0.5`) — and no more than two controls per treated unit (`max.controls = 2`). Under these restrictions, allowable matched sets could include 2 treated units with 1 control, 1 treated unit with 1 control (a matched pair), or 1 treated unit with 2 controls.

```

148 # Full matching using overall distance matrix
149 fullmatch(
150   x = overall_dist_mat,
151   min.controls = 0.5, # At least 0.5 controls per treated unit
152                      # (i.e., no more than 2 treated per control)
153   max.controls = 2,   # At most 2 controls per treated unit
154   omit.fraction = NULL, # Governs fraction of units discarded
155   mean.controls = NULL, # Governs average controls per treated unit
156   # Only one of omit.fraction or mean.controls can be non-NULL
157   data = data
158 )

```

Note, in addition, that the `fullmatch()` arguments of `omit.fraction` and `mean.controls` provide explicit levers for controlling how many treated and control units are discarded, though at most one of `omit.fraction` or `mean.controls` may be specified.

Imposing ratio constraints can have mixed consequences. In some cases, balance may worsen if a control is forced to match with a less similar treated unit — though still within the specified calipers — in order to satisfy the minimum and maximum ratio rules. In other cases, such constraints may improve effective sample size by redistributing how units are grouped without hurting balance. However, if the ratio constraints are too stringent, they can reduce effective sample size by forcing too many units to be discarded.

In applied settings, final choices of calipers and ratio constraints typically follow iterative checks of both covariate balance and effective sample size. Practitioners often compare several specifications and select among those that satisfy diagnostics for both goals. Hansen and Sales (2015) outline how this process can be carried out in a structured way, drawing on the stepwise intersection–union principle (SIUP) of hypothesis testing.

2.2.4 HOW DO I ACTUALLY FORM THE MATCHES?

The simple matching example above — based on an exact match on geographic region and a caliper on Euclidean distance for a single covariate (Polity score) — serves to illustrate the basic ideas. When matching on many covariates, however, we will often prefer some combination of Mahalanobis distance and propensity score matching, sometimes adding calipers on specific covariates. In what follows, we use *rank-based* Mahalanobis distance, which has the advantage of being less sensitive to outliers and differences in scales across covariates (Rosenbaum, 2010). We further constrain the matching by imposing a caliper on the estimated propensity score, requiring treated and control observations to come from the same geographic region, and applying additional calipers directly to two covariates: ethnic fractionalization (`ethfrac`) and logged GDP per capita (`bwgdp`).

We impose a caliper equal to 0.5 standard deviations of the logit of the estimated propensity score (the estimated linear covariate index defined above). This choice is larger than one rule of thumb emanating from Cochran and Rubin (1973), which recommends a caliper less than or equal to 0.20 standard deviations. Given that the standard deviation of the logit index is approximately 1.96, our choice of 0.5 permits treated and control units to differ by up to roughly 0.98 on the logit scale, compared to about 0.39 under the 0.20 guideline.

```

159 # Add logistic regression linear predictors to dataset
160 data$logit_p_score <- lin_cov_inds
161
162 # Population standard deviation of logit_p_score (divides by n, not n - 1)
163 pop_sd_logit <- sqrt(
164   mean(
165     (data$logit_p_score - mean(data$logit_p_score))^2
166   )
167 )
168
169 # Distance matrix for logit of estimated propensity scores
170 ps_mat <- match_on(
171   x = UN ~ logit_p_score,

```

```

172 caliper = 0.5 * pop_sd_logit,
173 data = data,
174 standardization.scale = NULL,
175 method = "euclidean"
176 )

```

Below we construct the distance matrix for rank-based Mahalanobis distance.

```

177 # Rank-based Mahalanobis distance on covariates
178 # Covs defined earlier as covariate names; here, drop "region"
179 rank_mah_mat <- match_on(
180   x = reformulate(
181     termlabels = setdiff(
182       x = covs,
183       y = "region"),
184     response = "UN"),
185   data = data,
186   standardization.scale = NULL,
187   method = "rank_mahalanobis" # Use rank-based Mahalanobis distance
188 )

```

Finally, we construct the Euclidean distance matrix for ethnic fractionalization (`ethfrac`) and logged GDP per capita (`bwgdp`) using calipers of 35 and 2, respectively. The exact-match constraint on region has already been defined through the object `em_region` above.

```

189 # Compute Euclidean distance matrix for ethnic fractionalization
190 eth_mat <- match_on(
191   x = UN ~ ethfrac,
192   caliper = 35,
193   data = data,
194   standardization.scale = NULL,
195   method = "euclidean"
196 )
197
198 # Compute Euclidean distance matrix for logged GDP per capita
199 bwgdp_mat <- match_on(
200   x = UN ~ bwgdp,
201   caliper = 2,
202   data = data,
203   standardization.scale = NULL,
204   method = "euclidean"
205 )

```

We then combine the `ps_mat`, `rank_mah_mat`, `eth_mat`, `bwgdp_mat`, and `em_region` objects to form the overall distance structure. We then pass the combined object to `optmatch`'s `fullmatch()` function, imposing a constraint that no more than 4 control units may be matched to any treated unit. If instead we wanted to perform pair matching, the `optmatch` package allows users to directly specify a matched-pair structure via the `pairmatch()` function. Equivalently, we could implement pair matching by setting `min.controls = 1` and `max.controls = 1` in the `fullmatch()` call. The code chunk below proceeds with full matching under the constraint that no more than 4 controls may be matched to any treated unit.

```

206 # Full matching:

```

```

207 # PS + rank Mahalanobis + Euclidean (ethfrac, bwgdp) + region exact
208 fm <- fullmatch(
209 # x specifies the distance structure:
210 # match_on() formula, distance matrix, or sum of match_on() distances
211 x = ps_mat + rank_mah_mat + eth_mat + bwgdp_mat + em_region,
212 data = data,
213 max.controls = 4 # <= 4 controls per treated (min.controls = 0 by default)
214 )

```

To calculate the effective sample size of the matched observations, we use the following function from `optmatch`.

```

215 # Effective sample size of matched sets
216 effectiveSampleSize(
217   fm
218 )

```

```
[1] 15.26667
```

Referring back to (2), this reported effective sample size of approximately 15.27 equals the sum across sets of the within-set harmonic mean of the number of treated and control subjects. This effective sample size reflects the matched structure in which no set contains more than 4 controls for any 1 treated observation.

```

219 # Summarize matched sets and report effective sample size
220 summary(
221   fm
222 )

```

```

Structure of matched sets:
1:0 2:1 1:1 1:2 1:3 1:4 0:1
  6  1  4  4  2  1 45
Effective Sample Size: 15.3
(equivalent number of matched pairs).

```

The notation in the `summary()` output of `optmatch` indicates the ratio of treated to control observations within each matched set. For example, 1:0 indicates 1 treated unit and no controls (effectively an unmatched treated unit). Similarly, 1:2 indicates 1 treated unit and 2 controls, and 0:1 indicates no treated units and 1 control (effectively an unmatched control). Below each label, the output shows how many matched sets have that particular structure.

If we examine the object returned by the matching call (`fm`), we see that `optmatch` labels each observation according to its matched set, assigning `NA` to those not included. Because we performed exact matching by region, `optmatch` labels each matched set using the name of the exact-match stratum followed by a set index within that stratum. For example, the label `lamerica.1` denotes the first matched set within the Latin America stratum.

To see which units were matched together, we can add the `fm` object to the dataframe and then tabulate. For example, to view the `ssafrica.3` set, we run the following.

```

223 # Add matched set ID to data for each unit

```

```

224 data$fm <- fm
225
226 # Look at one matched set ("ssafrica.3") for illustration
227 data |>
228   filter(
229     fm == "ssafrica.3" # Keep only units in set "ssafrica.3"
230   ) |>
231   # Display selected variables
232   select(
233     cname, UN, region, logit_p_score, ethfrac, bwgdp, bwplty2
234   )

```

```

# A tibble: 4 x 7
  cname      UN region  logit_p_score ethfrac bwgdp bwplty2
<chr>   <dbl> <fct>          <dbl>    <dbl> <dbl>   <dbl>
1 Burundi     0 ssafrica      0.508     3.55  5.35    -7
2 Rwanda      1 ssafrica      0.148    12.9  5.68    -7
3 Somalia     0 ssafrica      0.209     7.67  6.61    -7
4 Lesotho     0 ssafrica      0.558    22.2  6.28     0

```

We can see that the exact match on geographic region holds: All 4 countries are located in Sub-Saharan Africa. The logit of the estimated propensity score is similar across units, though not identical. The treated country, Rwanda, is matched to three controls — Burundi, Somalia, and Lesotho — and in each case the distance falls within 0.5 standard deviations of the estimated logit propensity scores (approximately 0.98), the caliper we specified. Likewise, Rwanda's distances to each of the 3 control units fall within the calipers of 35 for ethnic fractionalization (**ethfrac**) and 2 for logged GDP per capita (**wbgdp**). By contrast, distances between control countries can technically exceed these thresholds, since **optmatch** enforces calipers only between treated and control units, not among controls.

We can also see that some covariates used in the propensity score model and the rank-based Mahalanobis distance — such as Polity score (**wbplty2**) — still show modest imbalance within this matched set. This imbalance is unsurprising. We applied a caliper on the logit of the estimated propensity score and matched on the rank-based Mahalanobis distance including **wbplty2**, but we did not apply a caliper directly to **wbplty2**, though such a caliper could easily be added if desired.

2.2.5 HOW DO I DECIDE WHETHER TO MOVE AHEAD WITH MY MATCHED DESIGN?

Once we have constructed our matched sets, we want to evaluate the overall quality of the matched design. Covariate balance is an important aspect of this evaluation. To assess covariate balance, we compare the balance in our matched observational study with the balance we would expect to see in a completely randomized experiment within blocks (Hansen and Bowers, 2008).

Below we calculate adjusted covariate means for the treated and control groups by averaging set-specific means across matched sets, with each set weighted by its contribution to the effective sample size. We also report standardized differences, defined as the treated-minus-control difference in these adjusted means divided by the pooled standard deviation of the

covariate across treated and control units in the overall data (including both matched and unmatched units).

```

235 # Install "RIttools" package (only run if not already installed)
236 # install.packages("RIttools")
237
238 # Load RIttools package for balance diagnostics (balanceTest)
239 library(RIttools)
240 # Covariate balance test
241 cov_bal <- balanceTest(
242   # Formula: treatment ~ covariates
243   # update(): keep the original formula (. ~ .)
244   # and add stratification by matched set, + strata(fm)
245   fmla = update(
246     cov_fmla, . ~ . + strata(fm)
247   ),
248   data = data,
249   p.adjust.method = "none" # Method of p-value adjustment (none here)
250 )

```

Covariate	Before matching			After matching		
	Control mean	Treated mean	Std. diff	Control mean	Treated mean	Std. diff
Log Cumulative Battle Deaths from Last War	6.65	8.98	0.84*	8.34	8.57	0.08
Duration of Last War	50.28	80.53	0.39	60.51	73.23	0.16
Ethnic Fractionalization	56.50	49.21	-0.28	51.75	57.50	0.22
Log Population Size	9.51	8.75	-0.64*	8.89	8.91	0.02
Log Mountainous	2.22	2.80	0.43	2.69	2.82	0.09
Log Military Personnel	3.87	3.25	-0.42	3.63	3.54	-0.06
Log GDP per Capita Before Last War	6.56	6.59	0.03	6.73	6.55	-0.17
Democracy (Polity Score) Before Last War	-0.84	-2.58	-0.32	-2.19	-2.54	-0.07
Asia	0.19	0.00	-0.68*	0.00	0.00	0.00
Eastern Europe	0.15	0.37	0.51*	0.46	0.46	0.00
Latin America	0.12	0.21	0.25	0.08	0.08	0.00
North Africa & Middle East	0.12	0.11	-0.04	0.08	0.08	0.00
Sub-Saharan Africa	0.43	0.32	-0.23	0.38	0.38	0.00

Table 1: Adjusted covariate means for treated and control groups and standardized treated-control differences, before and after matching. Asterisks denote two-tailed p -values less than or equal to 0.05.

In Table 1 above, the adjusted means offer a direct description of balance. The standardized differences place all covariates on a common scale, enabling at-a-glance comparisons of imbalance across covariates. Stars indicate covariates for which the adjusted mean difference would be unusually extreme under a block-randomized experiment that assigns treatment completely at random within matched sets, holding fixed the observed number of treated units in each set. The stars are based on a Normal approximation to the randomization distribution of adjusted mean differences, and, because no adjustments are made for multiple comparisons, these stars are conservative. That is, for any given covariate, the corresponding adjusted mean difference is, if anything, more likely than the nominal rate to fall in the tails of its own randomization distribution under complete random assignment within matched sets. Alternative adjustments for testing multiple covariates are available, such as the Holm procedure (Holm, 1979), which is the default option in `RIttools`'s `balanceTest()` function.

One concern with balance tests is that high p -values may arise not from improved covariate balance but from the reduction in effective sample size that typically accompanies the matching process (Austin, 2008; Imai et al., 2008). As Hansen (2008b) notes, however, this possibility is less troubling than it first appears. The same increase in standard errors that produces high p -values for covariate balance tests will also carry over to subsequent causal inferences, meaning that those high p -values remain informative: They suggest that we are, if anything, less likely to overstate our causal conclusions than if the p -values had been significant.

Regardless of whether the balance tests are statistically significant, there are also established guidelines for what constitutes sufficient balance. While precise thresholds depend on context and substantive judgment about each covariate’s importance, two rules of thumb appear in Austin (2009) and Stuart (2010). Austin (2009) suggests that standardized differences of 0.1 or greater indicate inadequate balance on a covariate, whereas Stuart (2010), following Rubin (2001), proposes a more lenient threshold of 0.25. In our case, all covariates meet this latter standard, and none show statistically significant differences, indicating that observed imbalances would not be unusual under a completely randomized experiment within blocks (i.e., under as-if randomization).

In addition to assessing balance on each covariate individually, Hansen and Bowers (2008) also propose an omnibus, chi-squared (χ^2) test that evaluates balance across all covariates and their linear combinations simultaneously. We conduct this test below.

```
251 # Extract overall chi-square balance test by matched set (fm)
252 cov_bal$overall["fm", ]
```

```
      chisquare df    p.value
fm  2.631007   8 0.9553417
```

This χ^2 balance test yields an observed test statistic of 2.63 and a corresponding p -value of 0.96. Our high p -value indicates that the observed level of covariate balance is consistent with what we would expect in a completely randomized experiment within blocks.

Despite the high p -value, there is no guarantee that balance is sufficient for as-if randomization. Residual imbalance on observed covariates and hidden imbalance on unobserved ones may undermine the as-if randomization assumption. For now, we proceed under the as-if randomization assumption, but we will later assess how sensitive our inferences are to violations of it due to such imbalances.

2.3 Second Stage: Causal Inference (under As-If Randomization)

In both the sharp and weak frameworks, the causal targets of inference are defined through potential outcomes. Under what is known as SUTVA, the stable unit treatment value assumption (Rubin, 1980, 1986; Cox, 1958), each unit has two potential outcomes: a value the outcome would take if that unit were assigned to treatment and a value the outcome would take if that unit were assigned to control. Let $y_{si}(1)$ and $y_{si}(0)$ denote these potential outcomes, respectively, for unit i in set s , where the index i runs over the $\{1, \dots, n_s\}$ units in set s . The individual treatment effect is $\tau_{si} = y_{si}(1) - y_{si}(0)$. With $n = \sum_{s=1}^S n_s$ total

units, let $\boldsymbol{\tau} = (\tau_{1,1}, \tau_{1,2}, \dots, \tau_{S,n_s})^\top$ collect all n unit-level effects, indexed first by matched set and then by unit within set, and write the ATE as $\tau = (1/n) \sum_{s=1}^S \sum_{i=1}^{n_s} \tau_{si}$.

Both $\boldsymbol{\tau}$ and τ are defined only for the subset of units retained after matching, which often differs from the original set of units prior to matching (Rosenbaum and Rubin, 1985; Rosenbaum, 2012). This difference is especially important when retention is driven by overlap in the estimated propensity score, as the resulting study population can be difficult to interpret substantively. In such settings, it is often preferable to define the study population directly in terms of a small number of substantively meaningful covariates, reflecting the argument in Rosenbaum (2010) that it is “usually better to go back to the covariates themselves . . . perhaps redefining the population under study to be a subpopulation of the original population” (p. 86), as also emphasized by Stuart (2010). Building on this insight, subsequent approaches seek to define the matched study population directly in terms of a small number of substantively meaningful covariates (Fogarty et al., 2016; Traskin and Small, 2011).

The breadth of the study population retained after matching depends in part on the relative abundance of treated and control units across covariate space. When one treatment arm is sparsely represented in regions dominated by the other, overlap may be limited to a narrow region of covariate space and include relatively few units. All else equal, a larger pool of units available for matching supports overlap over a wider region of covariate space and a larger, more interpretable study population over which the target of inference is defined.

These targets of inference, $\boldsymbol{\tau}$ nor τ , defined among the matched units, cannot be directly calculated. Even in a randomized experiment, we cannot assign a country emerging from civil war to receive a UN peacekeeping mission, observe how long the ensuing peace lasts, and then rewind time to observe how long peace would have lasted in the absence of a UN peacekeeping mission. In other words, for each unit we observe only one of its two potential outcomes. Let y_{si} denote this observed outcome, which equals the treated potential outcome when unit i in set s receives the treatment and the control potential outcome otherwise. Because the individual-level causal effect can never be directly observed — only one potential outcome is realized for each unit — we must rely on statistical inference to draw conclusions about $\boldsymbol{\tau}$ and τ .

We consider inference under the sharp and weak frameworks, targeting $\boldsymbol{\tau}$ and τ , respectively. Ongoing work shows how both types of effects can be inferred simultaneously under as-if randomization (Chung and Romano, 2013; Ding, 2017; Wu and Ding, 2021; Cohen and Fogarty, 2022), though they cannot generally be unified in sensitivity analyses (Fogarty, 2023). When researchers must choose between the two, the decision depends on both statistical properties and substantive goals.

- Statistically, the sharp framework specifies all missing potential outcomes, allowing exact randomization inference under minimal assumptions. The weak framework leaves some outcomes unspecified and instead relies on variance estimation and a Normal approximation, which can perform poorly in small samples or when outcomes are skewed with extreme outliers. In such cases — or whenever one wants exact p -values

under minimal assumptions — permutation inference under the sharp framework may be preferable.

- Substantively, researchers usually test a constant effect in the sharp framework. Such a hypothesis may be unrealistic or of limited scientific interest (Gelman, 2003, 2011). The weak framework, by contrast, accommodates heterogeneous effects across units, making the ATE a more relevant target in many settings. Nevertheless, tests of a constant effect can provide a useful approximation to a more complex hypothesis with heterogeneous effects (Rosenbaum, 2010, pp. 44–46), and such tests remain valid for a range of bounded but heterogeneous effects (Caughey et al., 2023).

The Assignment Process as the Basis for Inference Regardless of the framework, inference is based on the treatment assignment process. Let z_{si} denote an indicator for whether unit i in matched set s is treated ($z_{si} = 1$) or not ($z_{si} = 0$). We collect these indicators for all units in set s into the vector $\mathbf{z}_s = (z_{s1}, \dots, z_{sn_s})^\top$. Stacking these vectors across all sets gives the full assignment vector $\mathbf{z} = (z_{11}, \dots, z_{Sn_S})^\top$, again indexed first by matched set and then by unit within set. For inference, we condition on the number of treated units within each set, m_s , even if the actual assignment mechanism were to consist of n_s independent assignments. For further discussion of this form of “conditional as-if analysis,” see Pashley et al. (2021) and Rosenbaum (2017, pp. 289–290, fn. 15).

We denote by Ω_s the set of all possible treatment assignments in set s , holding fixed the observed number of treated units in that set. Formally, Ω_s includes every possible way the n_s units in set s could be assigned to treatment and control such that exactly m_s units are treated. The number of possible assignments in Ω_s is denoted by $|\Omega_s|$, where the notation $|\cdot|$ indicates the number of elements in a set. This quantity equals $|\Omega_s| = \binom{n_s}{m_s} = n_s! / [m_s!(n_s - m_s)!]$, where “!” denotes the factorial operator (e.g., $4! = 4 \times 3 \times 2 \times 1$).

In the `ssafrica.3` set, for example, there are 4 observations — 1 treated and 3 control — so $|\Omega_s|$ for the `ssafrica.3` set is $\binom{4}{1} = 4$. The corresponding set of possible assignments with this treated count is shown in the table below.

	Assignment 1	Assignment 2	Assignment 3	Assignment 4
Burundi	0	0	0	1
Rwanda	0	0	1	0
Somalia	0	1	0	0
Lesotho	1	0	0	0

Table 2: All possible treatment assignments within matched set `ssafrica.3`, holding fixed the observed number of treated units in that set.

The column labeled Assignment 3 shows the assignment that actually occurred. The other possible assignments — Assignment 1, Assignment 2, and Assignment 4 — represent cases in which Lesotho, Somalia, or Burundi is treated instead of Rwanda. The set excludes any assignments with more than 1 treated unit.

The set of possible treatment assignments across all matched sets, given the number treated in each, is $\Omega = \Omega_1 \times \dots \times \Omega_S$, which is all the ways one assignment can be chosen from each Ω_s .

at the same time. Although the assignment itself can vary, the underlying causal quantities of interest — τ and τ — remain fixed across all possible assignments. What changes from one assignment to another is which potential outcomes we would actually observe. In the observed Assignment 3, we see Rwanda’s treated potential outcome and the control potential outcomes of Burundi, Somalia, and Lesotho, but not Rwanda’s control potential outcome or the treated potential outcomes of the others. Under a different assignment, a different set of treated and control potential outcomes would have been observed. No matter which assignment occurs, we observe only partial information about our causal targets.

Inference from the partial information contained in the data to our causal targets is predicated on a probability distribution defined over the set of assignments, Ω . This distribution constitutes the uncertainty underlying our inferences — what Fisher (1935, p. 14) famously called the “reasoned basis” for inference — but, unlike in a randomized experiment, this distribution is unknown in an observational study. Because this distribution is unknown, we must make assumptions about it when drawing inferences from observational data. Under the assumption of as-if randomization, every possible assignment within each Ω_s is equally likely, making each overall assignment in Ω equally likely as well. In this case, all individuals within the same matched set have the same probability of treatment, equal to m_s/n_s . When as-if randomization does not hold, however, assignments are no longer equally likely, and some individuals have a higher probability of treatment than others.

2.3.1 HOW DO I DRAW INFERENCES UNDER THE SHARP FRAMEWORK?

To set the stage for inference under the sharp framework, consider a thought experiment. Suppose we were to subtract the true individual effects from the outcomes of whichever units happen to be treated. Doing so would yield, for each unit, the outcome it would have had under control. In other words, we may imagine reconstructing the dataset so that, under any possible treatment assignment, the outcomes would appear exactly as they would have if no one had been treated. In this reconstructed world, there would be no effect since every unit’s outcome would reflect what it would have been without treatment.

Of course, we do not know the true collection of individual effects, τ , but we can test hypotheses about it, such as the hypothesis of a homogeneous effect for all units, denoted by τ_h . We do so by evaluating whether the data would look consistent with no treatment effect when that hypothesized value is subtracted from the treated outcomes. If, after this reconstruction, the data still show a positive effect, then the hypothesized value is presumably too small; if they show a negative effect, then the hypothesized value is presumably too large.

From a hypothesis testing perspective, when outcomes are reconstructed under a null hypothesis, the observed test statistic — computed on these reconstructed outcomes — will tend to fall in the upper tail of the null distribution if the hypothesized effect is too small, and in the lower tail if it is too large. To generate this null distribution, we reconstruct the outcomes under the hypothesized effect. Under the null, these reconstructed outcomes would remain fixed across assignments, so we hold them constant and recalculate the test statistic for every possible assignment. We then use this null distribution to determine where the observed test statistic falls for the calculation of p -values.

A canonical choice of test statistic in this setting is a *sum statistic*, which first adds up the treated outcomes within each set, and then adds those set-level sums across all sets. Many familiar test statistics can be expressed in this form by applying to the outcomes scale and shift transformations that do not depend on the treatment assignments. One such useful test statistic that can be expressed as a sum statistic is the Difference-in-Means, computed within sets and then averaged across sets, with each set’s contribution weighted by its effective sample size.

Under as-if randomization, the harmonic-mean-weighted Difference-in-Means is useful because, as discussed earlier, the variance of the within-set Difference-in-Means is minimized — under homogeneous treatment effects — when the harmonic mean of the numbers of treated and control units is largest. As a result, under this model of homogeneous individual treatment effects, a test has greater power when it places more weight on sets with large harmonic means — that is, sets in which the Difference-in-Means is most informative. Weighting sets solely by their share of units, by contrast, can overweight large but highly lopsided sets and dilute information from more balanced comparisons.

The harmonic-mean-weighted Difference-in-Means also has the practical advantage of coinciding with two other common approaches to analyzing matched data:

- First, the harmonic-mean-weighted Difference-in-Means equals the coefficient on the treatment indicator from a fixed-effects (FE) regression that includes matched set indicators (Hansen and Bowers, 2008, pp. 228-229), an approach commonly used in practice after matching.
- Second, the harmonic-mean-weighted Difference-in-Means coincides with the overlap-weighted Difference-in-Means (Li et al., 2018) when overlap weights are constructed using the treatment assignment probabilities implied by as-if randomization. Overlap weights place the greatest weight on units with intermediate values of these assignment probabilities and downweight units with extreme values. In particular, using the assignment probability under as-if randomization, control units are weighted by that probability, while treated units are weighted by 1 minus that probability. Because these assignment probabilities are constant within each matched set, the resulting overlap weights are constant within treated units and within control units in a set, and they aggregate exactly to the harmonic-mean weights implied by the matched sets.

To implement a sum statistic equivalent to the harmonic-mean-weighted Difference-in-Means for hypothesis testing, we first reconstruct the outcomes that treated units would have exhibited under the null hypothesis $\tau_h = 0$.

```

253 # Keep only rows assigned to a matched set (drop NA in fm)
254 data_matched <- filter(
255   .data = data, !is.na(fm)
256 )
257
258 # Null hypothesis value
259 tau_h <- 0
260
261 # Reconstruct outcomes under sharp null (tau_h = 0)
262 data_matched <- data_matched |>

```

```

263 mutate(
264   ldur_tilde = ldur - tau_h * UN
265 )

```

We now source a helper function from the companion GitHub repo; this function rescales the outcome variable so that the sum of the rescaled values among treated units equals the harmonic-mean-weighted Difference-in-Means. After sourcing the function, we apply it to the matched dataset to generate a new column containing the rescaled outcomes. We then use this rescaled outcome to compute the observed test statistic.

```

266 # Load hm_stat_rescale() from GitHub repo
267 # base_url (defined earlier) points to repo URL
268 source(
269   paste0(
270     base_url, "/R/hm_stat_rescale.R"
271   )
272 )
273
274 # Apply rescaling: adds .hm_scaled and returns matched dataset
275 data_matched <- hm_stat_rescale(
276   data = data_matched,
277   outcome = ldur_tilde, # Set outcome to be rescaled within matched sets
278   treat = UN,           # Set name of treatment indicator variable
279   strata = fm           # Set name of matched strata (block) variable
280 )
281
282 # Observed HM-weighted diff-in-means statistic
283 obs_stat <- sum(
284   data_matched$ldur_tilde_hm_scaled[data_matched$UN == 1]
285 )

```

We can verify in R that this observed sum statistic coincides with the harmonic-mean-weighted Difference-in-Means, the coefficient on the treatment indicator from a FE regression that includes matched set indicators, and the overlap-weighted Difference-in-Means constructed from the assignment probabilities implied by as-if randomization.

First, we show that the observed sum statistic equals the harmonic-mean-weighted average of the within-matched-set Differences-in-Means, where each matched set is weighted by its contribution to the effective sample size.

```

286 # Harmonic-mean-weighted difference in means (weights computed within sets)
287 dim_hm <- data_matched |>
288   group_by(
289     fm
290   ) |>
291   summarize(
292     n_treated = sum(UN == 1),
293     n_control = sum(UN == 0),
294     # Within-set Difference-in-Means
295     dim_set   = mean(ldur_tilde[UN == 1]) - mean(ldur_tilde[UN == 0]),
296     # Harmonic-mean weight (set's contribution to effective sample size)
297     w_hm      = 2 * n_treated * n_control / (n_treated + n_control),
298     .groups   = "drop" # Drop grouping after summarise
299   ) |>

```

```

300   summarize(
301     # Harmonic-mean-weighted average of within-set Differences-in-Means
302     dim_hm = sum(
303       w_hm * dim_set) / sum(w_hm)
304   ) |>
305   pull( # Pull column out of data frame
306     dim_hm
307   )
308
309 # Approaches coincides with the sum statistic
310 all.equal(
311   dim_hm, obs_stat
312 )

```

[1] TRUE

Next, we verify that the same quantity is obtained as the coefficient on the treatment indicator from a FE regression that includes matched set indicators.

```

313 # FE coefficient on UN (matched-set FE)
314 fe_fit <- lm(
315   formula = ldur_tilde ~ UN + fm,
316   data = data_matched
317 )
318
319 # Drop name so comparison is purely numeric
320 treat_coef_fe_fit <- unname(
321   coef(fe_fit)["UN"]
322 )
323
324 # Approaches coincides with the sum statistic
325 all.equal(
326   treat_coef_fe_fit, obs_stat
327 )

```

[1] TRUE

Finally, we confirm that the observed sum statistic also coincides with the overlap-weighted Difference-in-Means. The overlap weights are constructed from the treatment assignment probabilities implied by as-if randomization within matched sets, with treated units weighted by one minus their treatment probability and control units weighted by their treatment probability.

```

328 # Install "PSweight" package (only run if not already installed)
329 # install.packages("PSweight")
330 library(PSweight) # For overlap weights from Li et al (2018)
331
332 # Treatment probabilities within matched sets: n_treated / n
333 ps_fit <- PSmethod(
334   ps.formula = UN ~ factor(fm), # Reproduces n_treated / n in each set
335   method     = "glm",           # Estimate using generalized linear model
336   # Use as.data.frame() for PSmethod() compatibility
337   data       = as.data.frame(data_matched),
338   ncate      = 2                 # Binary treatment (treated vs. control)
339 )

```

```

340
341 assign_prob <- ps_fit$e.h[, "1"] # Column for treatment level "1"
342
343 # Overlap weights from treatment probabilities
344 # Controls: treatment probability ; Treated: 1 - treatment probability
345 w_ow <- ifelse(
346   test = data_matched$UN == 1,
347   yes = 1 - assign_prob,
348   no = assign_prob
349 )
350
351 # Overlap-weighted difference in means
352 dim_ow <- data_matched |>
353   summarize( # Aggregate rows into summary values
354     # Overlap-weighted mean outcome for treated units
355     treated_weighted_mean =
356       sum(w_ow[UN == 1] * ldur_tilde[UN == 1]) / sum(w_ow[UN == 1]),
357
358     # Overlap-weighted mean outcome for control units
359     control_weighted_mean =
360       sum(w_ow[UN == 0] * ldur_tilde[UN == 0]) / sum(w_ow[UN == 0]),
361
362     # Difference between overlap-weighted treated and control means
363     dim_ow = treated_weighted_mean - control_weighted_mean
364   ) |>
365   pull(
366     dim_ow
367   )
368
369 # Approaches coincides with the sum statistic
370 all.equal(
371   dim_ow, obs_stat
372 )

```

[1] TRUE

Now, to generate the distribution to which we refer our observed sum statistic, we can hold the reconstructed and rescaled outcomes fixed and then calculate the sum statistic over all possible assignments in Ω . In our application, with 12 sets ranging in size from 2 to 5, the total number of assignments is 311,040.

```

373 # For each matched set (fm), record:
374 # n = total units in the set
375 # m = number treated (UN == 1)
376 block_ns <- data_matched |>
377   group_by( # Group results by key variables
378     fm
379   ) |>
380   summarise( # Aggregate to one row per group
381     n = n(), # Row count per group
382     m = sum(UN),
383     .groups = "drop"
384   )
385
386 # Total possible treatment assignments = product of binomial coefficients

```

```

387 # (choose n_s units for treatment in each set and multiply across sets)
388 exact_n_assigns <- prod(
389   choose(
390     n = block_ns$n,
391     k = block_ns$m
392   )
393 )

```

Because our matched study is relatively small, we can enumerate all possible treatment assignments exactly.

```

394 # Install "randomizr" (only run if not already installed)
395 # install.packages("randomizr")
396
397 # Load randomizr for generating random assignments
398 library(randomizr)
399
400 exact_assigns <- obtain_permutation_matrix(
401   declaration = declare_ra( # Declare assignment procedure
402     N = nrow(data_matched), # Total number of units
403     blocks = data_matched$fm, # Matched set membership
404     block_m = block_ns$m      # Number treated in each set
405   ),
406 # Total number of assignments with treated counts fixed within sets
407 maximum_permutations = exact_n_assigns
408 )

```

However, in most applications, exactly enumerating all possible assignments is computationally infeasible. Instead, we typically draw a random subset of assignments (e.g., 10,000) to approximate the exact randomization distribution, as illustrated below.

```

409 # Set RNG seed for reproducibility
410 set.seed(11242017)
411
412 # Generate permutation matrix of treatment assignments
413 # Each column = one possible assignment consistent with block structure
414 sim_assigns <- obtain_permutation_matrix(
415   declaration = declare_ra(
416     N = nrow(data_matched),
417     blocks = data_matched$fm,
418     block_m = block_ns$m
419   ),
420   maximum_permutations = 10^4 # Cap at 10,000 random draws
421 )

```

We apply the statistic to each of these 10,000 assignments while holding the reconstructed and rescaled outcomes fixed under the null.

```

422 # Randomization distribution under sharp null
423 # Apply sum statistic to each assignment in sim_assigns
424 # Outcome transformed so sum = harmonic-mean weighted diff in means
425 sim_sharp_null_dist <- apply(
426   X = sim_assigns, # Matrix of treatment assignments
427   MARGIN = 2,      # Iterate over columns (assignments)
428   FUN = function(x) {

```

```

429 # Sum transformed outcomes among treated
430 sum(
431   data_matched$ldur_tilde_hm_scaled[x == 1]
432 )
433 }
434 )
435 # Faster equivalent computation via matrix multiplication
436 # as.numeric(
437 #   t(data_matched$ldur_tilde_hm_scaled) %*% sim_assigns
438 # )

```

From this null distribution, we compute a one-sided upper p -value, which, under as-if randomization, is simply the proportion of null test statistics greater than or equal to the observed test statistic.

```

439 # One-sided, upper p-value:
440 # Propprtion of simulated randomization stats >= observed
441 round(
442   x = mean(
443     sim_sharp_null_dist >= obs_stat
444   ),
445   digits = 4
446 )

```

```
[1] 0.0346
```

The upper p -value of a test of the sharp null of no effect against the alternative of a larger effect is 0.0346.

In this particular case, unlike in most applications, the matched study is small enough to compute the exact p -value and assess the accuracy of the simulation-based p -value approximation.

```

447 # Transformed outcomes under sharp null (tau_h = 0)
448 q_tau_h_0 <- data_matched$ldur_tilde_hm_scaled
449
450 # Fast computation of exact null distribution using matrix multiplication
451 exact_sharp_null_dist <- as.numeric(
452   t(q_tau_h_0) %*% exact_assigns
453 )
454
455 # Slower apply()-based computation (for reference)
456 # apply(
457 #   X = exact_assigns,
458 #   MARGIN = 2,
459 #   FUN = function(x) {
460 #     sum(data_matched$ldur_tilde_hm_scaled[x == 1])
461 #   }
462 # )
463
464 # Exact one-sided, upper p-value
465 round(
466   x = mean(
467     exact_sharp_null_dist >= obs_stat
468   ),

```



```

469   digits = 4
470 )

```

```
[1] 0.0363
```

The exact p -value of 0.0363 is nearly identical to the simulation-based p -value of 0.0346. At the conventional significance level of $\alpha = 0.05$, we would reject the null hypothesis in favor of the alternative, regardless of whether we use the exact or simulation-based p -value.

As an alternative to randomly sampling assignments and computing the test statistic for each one, we can use a much faster Normal approximation to the null distribution when the matched study is sufficiently large. The approximation relies on closed-form expressions for the expected value and variance of a sum statistic derived in Rosenbaum and Krieger (1990). Using these expressions, we standardize the observed test statistic and then compare it to the standard Normal distribution, which gives us a corresponding p -value. Although primarily designed for sensitivity analyses under violations of as-if randomization, the `senstrat` package can also be used to compute the null expected value and variance of a sum statistic under as-if randomization.

```

471 # Install "senstrat" package (only run if not already installed)
472 # install.packages("senstrat")
473
474 # Load senstrat for stratum-level null moments and sensitivity analysis
475 # (Rosenbaum & Krieger 1990)
476 library(senstrat)
477
478 # Compute per-block null expectations and variances
479 per_block_moms <- data_matched |>
480   group_by(
481     fm
482   ) |>
483   summarize(
484     expect = ev(
485       sc = ldur_tilde_hm_scaled, # Transformed outcomes for the stratum
486       z = UN,                   # Treatment indicator
487       m = 1,                   # Count of 1s in hidden confounder
488                                # Irrelevant here since Gamma = 1
489       g = 1,                   # Sensitivity parameter Gamma
490       method = "RK"            # Use Rosenbaum and Krieger (1990) formula
491     )$expect,                  # Null expectation of sum statistic by set
492     variance = ev(
493       sc = ldur_tilde_hm_scaled,
494       z = UN,
495       m = 1,
496       g = 1,
497       method = "RK"
498     )$vari,                    # Null variance of sum statistic by set
499     .groups = "drop"
500   )
501
502 # Sum across blocks to get overall null expectation
503 null_ev <- sum(
504   per_block_moms$expect
505 )

```

```

506
507 # Sum across blocks to get overall null variance
508 null_var <- sum(
509   per_block_moms$variance
510 )
511
512 # Standardized test statistic and one-sided Normal p-value (upper tail)
513 norm_upper_p_value <- pnorm(
514   q = (obs_stat - null_ev) / sqrt(null_var), # Standardized statistic
515   lower.tail = FALSE                        # Compute upper-tail prob
516 )
517 # By default in pnorm(): mean = 0 and sd = 1 (standard normal distribution)

```

This Normal-approximation p -value of 0.0343 is close to the exact and simulation-based p -values of 0.0363 and 0.0346, respectively. All p -values lead to the same conclusion in which we reject the sharp null of no effect in favor of a larger effect.

To test not just a single null hypothesis but a grid of null hypotheses under the sharp null framework, we use a Normal approximation to the randomization distribution. We repeat the above procedure for a range of hypothesized constant effects, τ_h . For each value of τ_h , we reconstruct the outcomes under the null, rescale them so that the sum statistic equals the harmonic-mean-weighted Difference-in-Means, compute the observed test statistic, and obtain the corresponding p -value using the Normal approximation. To form a one-sided confidence set, we retain all null values that are not rejected by the upper-tail test at the chosen α level. To form a two-sided confidence set, we proceed analogously, allocating $\alpha/2$ to each tail — reflecting that each null value is tested against departures in both directions — and using the corresponding upper- and lower-tail rejection regions.

```

518 # Significance level
519 alpha <- 0.05
520
521 # Upper-tail confidence set (tau_h values not rejected by upper-tail test)
522 # Lower endpoint: smallest tau_h with upper-tail p-value >= alpha
523 # (any smaller tau_h would be rejected)
524 cs_sharp_lower_one_sided <- c(
525   lower = obs_stat - qnorm(p = 1 - alpha) * sqrt(null_var),
526   upper = Inf
527 )
528 # qnorm(1 - alpha): upper one-sided Normal critical value
529
530 cs_sharp_two_sided <- c(
531   lower = obs_stat - qnorm(p = 1 - alpha / 2) * sqrt(null_var),
532   upper = obs_stat + qnorm(p = 1 - alpha / 2) * sqrt(null_var)
533 )
534 # At lower endpoint, upper-tail one-sided p-value equals alpha/2;
535 # any smaller null value would be rejected by two-sided test
536 # At upper endpoint, lower-tail one-sided p-value equals alpha/2;
537 # any larger null value would be rejected by two-sided test

```

An analogous procedure applies when constructing confidence sets without relying on a Normal approximation, instead using a simulation-based approximation to the null randomization distribution.

```

538 # Grid of constant-effect null values (sharp framework)
539 tau_h_grid <- seq(
540   from = -0.02,
541   to = 1.5,
542   by = 0.0001
543 )
544
545 # Repeat the randomization test for each tau_h
546 # sapply() computes and stacks the resulting p-values
547 p_mat <- sapply(X = tau_h_grid,
548   FUN = function(tau_h) {
549
550     # Shift outcomes under null: ldur_i - tau_h * UN_i
551     dat_tau_h <- hm_stat_rescale(
552       # transform(): copy data_matched adding shifted outcome
553       data = transform(
554         data_matched,
555         ldur_tilde_shift = ldur - tau_h * UN),
556       outcome = ldur_tilde_shift,
557       treat = UN,
558       strata = fm
559     )
560
561     # Transformed outcomes under sharp null tau_h
562     q_tau_h <- dat_tau_h$ldur_tilde_shift_hm_scaled
563
564     # Observed statistic under null tau_h
565     obs_stat_tau_h <- sum(
566       dat_tau_h$UN * q_tau_h
567     )
568
569     # Randomization distribution via simulated assignments
570     # Matrix multiplication used for speed (vs apply())
571     sim_null_dist_tau_h <- as.numeric(
572       t(q_tau_h) %*% sim_assigns
573     )
574
575     # Upper-tail and lower-tail randomization p-values
576     p_upper_tau_h <- mean(
577       sim_null_dist_tau_h >= obs_stat_tau_h
578     )
579
580     p_lower_tau_h <- mean(
581       sim_null_dist_tau_h <= obs_stat_tau_h
582     )
583
584     c(
585       upper_tail = p_upper_tau_h,
586       lower_tail = p_lower_tau_h
587     )
588   })
589
590
591 # Extract vector of upper-tail p-values
592 p_upper <- p_mat["upper_tail", ]

```

```

593
594 # Extract vector of lower-tail p-values
595 p_lower <- p_mat["lower_tail", ]
596
597 # Simulation-based confidence sets
598
599 # Upper-tail confidence set: retain all tau_h with upper-tail p >= alpha
600 cs_sharp_upper_tail_sim <- tau_h_grid[p_upper >= alpha]
601
602 # Lower bound of the upper-tail confidence set
603 cs_sharp_upper_tail_sim_bound <- min(
604   cs_sharp_upper_tail_sim
605 )
606
607 # Two-sided confidence set (inversion using alpha/2 in each tail):
608 # retain tau_h only if neither one-sided test rejects at level alpha/2
609 cs_sharp_two_sided_sim <- tau_h_grid[
610   p_upper >= alpha / 2 & p_lower >= alpha / 2
611 ]
612
613 # Two-sided confidence set summarized by its bounds
614 cs_sharp_two_sided_sim_bounds <- c(
615   lower = min(
616     cs_sharp_two_sided_sim
617   ),
618   upper = max(
619     cs_sharp_two_sided_sim
620   )
621 )

```

Likewise, for a point estimate, one could follow Hodges Jr. and Lehmann (1963) and Rosenbaum (1993) by identifying the null value that makes the observed sum statistic equal to its null expectation. With our test statistic under as-if randomization, this would occur when the null equals the harmonic-mean weighted Difference-in-Means computed from the observed outcomes (roughly 0.673), which yields a null expectation of 0.

2.3.2 HOW DO I DRAW INFERENCES UNDER THE WEAK FRAMEWORK?

When the target is the ATE, it can be written as $\tau = \sum_{s=1}^S (n_s/n) \tau_s$, a weighted average of the set-level ATEs with weights equal to each set's share of the total study size. Thus, a straightforward way to estimate the ATE is to compute the Difference-in-Means within each matched set, and then combine these set-level estimates using the same weights. Formally, the overall Difference-in-Means is $\hat{\tau} = \sum_{s=1}^S (n_s/n) \hat{\tau}_s$, where $\hat{\tau}_s$ is the Difference-in-Means in set s .

These weights are appealing because, under as-if randomization, the resulting Difference-in-Means is an unbiased estimator of the ATE. Formally, the expected value of the estimator — i.e., the average of $\hat{\tau}$ taken over all treatment assignments consistent with the matched design and their associated probabilities — equals the true ATE. Nevertheless, although correct in expectation, the Difference-in-Means can vary substantially across assignments. In any given assignment, the estimate may lie far from the target ATE.

Because the estimator can fluctuate across different treatment assignments, it is important to quantify the typical squared distance between an estimate and the true ATE. To this end, Neyman (1923) introduced a canonical variance estimator. Under as-if randomization, this estimator is exactly unbiased when individual treatment effects are homogeneous; otherwise, it is conservative, meaning its expected value is greater than or equal to the true variance of the Difference-in-Means, although subsequent refinements reduce this conservativeness while maintaining validity (see Robins, 1988; Aronow et al., 2014; Harshaw et al., 2024). In principle, we could apply this approach by estimating the variance of the Difference-in-Means within each matched set and then taking a weighted sum of these estimates (see, e.g., Miratrix et al., 2013).

In our setting, however, this approach is infeasible because each matched set contains only 1 treated or 1 control unit. The usual Neyman formula relies on computing sample variances separately within the treated and control groups of each set. The sample variance requires at least two observations because `var()` divides by the number of observations minus 1. As a result, we cannot compute the sample variance when there is only a single treated or control unit.

How, then, can we estimate the variance of the Difference-in-Means in a matched observational study? Pashley and Miratrix (2021) and Fogarty (2018) offer distinct solutions. Pashley and Miratrix (2021) propose two approaches, each valid under different conditions on the matched structure, while Fogarty (2018) develops a single method that applies more broadly to finely stratified designs.

The first approach in Pashley and Miratrix (2021), `hybrid_m`, requires at least two matched sets of each unique set size in the data. The second approach, `hybrid_p`, permits variation in set sizes as long as no single set contains half or more of the total study size. Our matched design meets the condition for the second approach, not the first. We therefore estimate the variance using `hybrid_p` through the `blkvar` package, the companion software for Pashley and Miratrix (2021).

```
622 # Install blkvar (only run if not already installed)
623 # install.packages("remotes")
624 # Remotes::install_github("lmiratrix/blkvar")
625
626 # Load blkvar for block randomization variance estimators
627 library(blkvar)
628
629 # Compute results with hybrid_p method
630 res <- block_estimator(
631   Yobs = ldur,          # Observed outcomes
632   Z = UN,               # Treatment indicator
633   B = fm,               # Block (matched set) membership
634   data = data_matched, # Dataset
635   method = "hybrid_p"  # variance estimation method
636 )
637
638 # Extract variance estimate
639 res$var_est
```

```
[1] 0.1187282
```

We can also estimate the variance of the Difference-in-Means in a finely stratified design using the approach of Fogarty (2018). To do so, we load a custom R function, `fine_strat_var_est()`, from the companion GitHub repo.

```

640 # Load the fine_strat_var_est() function from the GitHub repo
641 source(
642   paste0(
643     base_url, "/R/fine_strat_var_est.R"
644   )
645 )

```

The function takes as inputs the set-specific treatment effect estimates and the number of units in each set, and it returns a single scalar representing the estimated variance. We compute the set-specific estimates and corresponding set sizes, and then pass these two vectors as inputs to `fine_strat_var_est()`.

```

646 # Compute matched set sizes and matched-set-specific differences in means
647 set_stats <- data_matched |>
648   group_by(
649     fm
650   ) |>
651   summarize(
652     n = n(), # Matched set size
653     diff_in_means = mean( # Treated mean
654       ldur[UN == 1L]
655     ) - # minus
656     mean( # control mean
657       ldur[UN == 0L]
658     ),
659     .groups = "drop"
660   )
661
662 # Apply Fogarty (2018/2023) variance estimator
663 fine_strat_var_est(
664   strat_ns = set_stats$n, # Vector of stratum sizes
665   strat_ests = set_stats$diff_in_means # Vector of stratum estimates
666 )

```

```
[1] 0.1163934
```

This variance estimate is nearly identical to the one we obtained using the `hybrid_p` approach of Pashley and Miratrix (2021).

Now that we have estimates of both the ATE and the variance of the ATE estimator, we can form a standardized test statistic by subtracting the expected Difference-in-Means implied by the null and dividing the result by the estimated standard error (the square root of the estimated variance). We then compare this statistic to a standard Normal distribution to calculate p -values. Below we calculate the upper one-sided p -value for a test the null hypothesis that the ATE is 0 against the alternative that it is positive.

```

667 pnorm(
668   q = (res$ATE_hat - 0) / sqrt(res$var_est),
669   lower.tail = FALSE
670 )

```

[1] 0.03722157

Here, we reject the weak null hypothesis that the ATE equals 0 in favor of the alternative that the ATE is greater than 0, using the conventional significance level of $\alpha = 0.05$.

As in the construction of confidence sets for a homogeneous treatment effect under the sharp framework, we form 95% confidence sets by inverting the corresponding hypothesis tests. We first construct a one-sided confidence set obtained by inverting the upper-tail test and retaining all null values that are not rejected at the $\alpha = 0.05$ level. We then construct a two-sided confidence set by allocating $\alpha/2 = 0.025$ to each tail, as in the tail allocation used in the sharp null framework.

```

671 # One-sided (upper-tail) confidence set
672 cs_weak_upper_tail <- c(
673   lower = res$ATE_hat - qnorm(p = 1 - alpha) * sqrt(res$var_est),
674   upper = Inf
675 )
676
677 # Two-sided confidence set
678 cs_weak_two_sided <- c(
679   lower = res$ATE_hat - qnorm(p = 1 - alpha / 2) * sqrt(res$var_est),
680   upper = res$ATE_hat + qnorm(p = 1 - alpha / 2) * sqrt(res$var_est)
681 )

```

2.4 Third Stage: Sensitivity Analysis (Departures from As-If Randomization)

Up to this point, we have supposed a framework in which each observation's chance of a UN intervention is like flipping a weighted coin. Each coin flip is independent across observations, but the probability of landing tails (i.e., receiving treatment) can differ from one unit to another depending on its baseline characteristics. With matching, we make our crucial assumption that all units within a set are similar enough on these characteristics that their coins have the same probability of landing tails.

This as-if randomization assumption may fail because of imbalances on hidden covariates (or residual imbalances on observed covariates, though the sensitivity analysis to follow subsumes both within the same framework). To represent such imbalances, consider a single hidden covariate, $\mathbf{u} = (u_{11}, \dots, u_{S n_S})^\top$, indexed first by matched set and then by unit within set, with each u_{si} constrained to lie in the interval from 0 to 1. Although the restriction that each element of $\mathbf{u}$ is between 0 and 1 may seem strong, Rosenbaum (2017, p. 300, fn. 33) shows that any departure from complete random assignment within blocks can be represented by such a $\mathbf{u}$ in the sense that as-if randomization would hold if we had access to and exactly matched on this $\mathbf{u}$.

Rosenbaum (1987a) and Rosenbaum and Krieger (1990) then propose a model in which each unit's independent probability of treatment assignment is given by

$$\pi_{si} = \frac{\exp[\kappa_s + \log(\Gamma)u_{si}]}{1 + \exp[\kappa_s + \log(\Gamma)u_{si}]} \quad (3)$$

The parameter κ_s is a set-specific intercept that captures the baseline propensity for treatment shared by all units in matched set s before accounting for any differences in the hidden

covariate. In other words, κ_s reflects the central idea in matching that, after forming matched sets homogeneous in observed covariates, all individuals within a set share the same treatment propensity based on those covariates. The parameter $\Gamma \geq 1$, by contrast, quantifies how strongly the hidden covariate u_{si} can alter treatment odds. When $\Gamma = 1$, all individuals in matched set s have the same probability of treatment, $\exp[\kappa_s]/(1 + \exp[\kappa_s])$. However, when $\Gamma > 1$, two individuals in the same set who differ in the hidden covariate may differ in their odds of treatment by as much as a factor of Γ .

An important point to reiterate is that we condition on assignments that belong to the set Ω , meaning that the number of treated units is fixed within each matched set. It turns out that this conditioning removes dependence on the set-specific baseline κ_s in the probability distribution over assignments in Ω (Rosenbaum, 1984). Eliminating this dependence is crucial because it allows us to characterize both as-if randomization and departures from it using just a single sensitivity parameter, Γ .

To build intuition for the sensitivity parameter Γ , note that the model in (3) implies the following restriction:

$$\frac{1}{\Gamma} \leq \frac{\pi_{si}(1 - \pi_{sj})}{\pi_{sj}(1 - \pi_{si})} \leq \Gamma \text{ for all } i, j \text{ and } s. \quad (4)$$

This restriction states that, within any set, no two units can differ in their odds of treatment by more than a factor of Γ . When $\Gamma = 1$, all units share the same odds of treatment, corresponding to as-if randomization. By contrast, larger values of Γ represent increasingly severe departures from this assumption.

Rosenbaum (1995, pp. 1424–1425) shows that the converse also holds: The restriction in (4) implies a model of the form in (3). For any collection of treatment probabilities satisfying the bound in (4), it is always possible to find values of $u_{si} \in [0, 1]$ for all sets s and units i , along with a scalar $\Gamma \geq 1$, such that the two formulations yield the same probability distribution over assignments in Ω . In other words, (4) describes a general restriction on treatment probabilities that does not assume a particular functional form. The logistic model in (3), by contrast, provides one convenient parametric representation of that general restriction.

Thus far, we have conducted inference assuming $\Gamma = 1$, meaning that all units within a matched set share the same treatment probability. Under this as-if randomization assumption, tests of sharp null hypotheses are exactly valid, while tests of weak null hypotheses are valid in large studies. In both cases, validity means that the probability of falsely rejecting the null hypothesis does not exceed the nominal level α .

When we allow departures from as-if randomization, governed by the sensitivity parameter $\Gamma \geq 1$, we seek new p -values that remain valid in the same sense. The probability of a false rejection should not exceed α , regardless of the hidden covariate $\mathbf{u}$ or the potential outcomes consistent with the null. The way we ensure this validity, however, differs between the sharp and weak frameworks, which we consider in turn.

2.4.1 HOW DO INFERENCES UNDER THE SHARP FRAMEWORK CHANGE UNDER THESE DEPARTURES?

For a sharp null, if the hidden $\mathbf{u}$ were known, we could calculate from (3) the exact distribution of assignments consistent with the matched design for any given value of $\Gamma \geq 1$. This distribution would provide the correct reference for computing p -values under the null. In particular, the upper one-sided p -value could be obtained by summing the probabilities of all assignments whose test statistic under the null is at least as large as the observed statistic.

In practice, however, $\mathbf{u}$ is unknown. Thus, to ensure validity of our tests, we compute p -values under the worst-case choice of $\mathbf{u}$ — the one that makes the p -value as large as possible. Rejection under the worst-case choice of $\mathbf{u}$ guarantees rejection under the actual (but hidden) $\mathbf{u}$. Consequently, we will not reject the null more often than we would if the true assignment probabilities — governed by the actual but unobserved $\mathbf{u}$ — were known under the model in (3).

Finding the Worst-Case Scenario of Hidden Confounding for Valid Inference

Actually finding this worst-case $\mathbf{u}$ is difficult. As a first step, Rosenbaum and Krieger (1990) show that in the unmatched (i.e., two sample) case, the vector $\mathbf{u}$ that maximizes the p -value under any fixed $\Gamma \geq 1$ must belong to a restricted set of possibilities, denoted by $\mathcal{U}_+$. Once the subjects are ordered from the largest to the smallest outcome, all the candidate vectors in $\mathcal{U}_+$ look the same: a sequence of 1s at the top followed by 0s below. We do not know *how many* 1s should precede the 0s, and this number determines which configuration of $\mathbf{u}$ yields the worst-case p -value. In an unmatched study, it is straightforward to enumerate all $n - 1$ such candidate vectors in $\mathcal{U}_+$, and then identify which one yields the largest p -value for a fixed $\Gamma \geq 1$.

Unfortunately, in matched designs, the overall worst-case vector $\mathbf{u}$ cannot be obtained by simply stitching together the worst-case vectors from each matched set. That is, the global worst-case is not just the collection of local worst-cases. Instead, $\mathcal{U}_+$ consists of $\prod_{s=1}^S (n_s - 1)$ total candidate vectors for the worst-case $\mathbf{u}$ — a quantity that quickly becomes infeasible to enumerate directly. For example, with only 20 matched sets of 4 units each, the number of candidate vectors already exceeds 3.5 billion.

Separable Approximation To address this challenge, Gastwirth et al. (2000) propose a practical shortcut called the *separable approximation*. The idea is to select, within each matched set, the configuration $\mathbf{u}_s$ (from the set-specific candidate sequences of 1s and 0s) that maximizes the expected value of the test statistic in that set under the null. If more than one candidate yields the same expectation, the choice goes to the $\mathbf{u}_s$ that produces the larger variance of the test statistic in that set under the null. The method is called “separable” because it then stitches together the choices of $\mathbf{u}_s$ made separately within each matched set, rather than searching over all possible combinations across sets. The resulting $\mathbf{u}$ may not give the exact worst-case p -value; yet in designs with many small matched sets, the error is negligible.

Taylor Series Approximation The separable approximation is a useful shortcut, but it can fail in designs with relatively few matched sets or with highly unbalanced set sizes. In such cases, the $\mathbf{u}$ selected by the separable approximation may yield p -values that are smaller than the true worst-case p -value. Rosenbaum (2018) therefore introduces a refinement that guarantees valid inference.

The key idea is to reframe hypothesis testing in terms of a function that, for each configuration of hidden confounding $\mathbf{u} \in \mathcal{U}_+$, combines the gap between the null expectation and the observed test statistic with a corresponding Normal critical buffer. Concretely, for $\alpha = 0.05$ and corresponding critical value 1.64, this function is the expected value of the test statistic under $(\mathbf{u}, \Gamma)$ minus the observed test statistic, plus 1.64 times the standard deviation of the test statistic under $(\mathbf{u}, \Gamma)$. The function is nonpositive exactly when the observed test statistic lies at least 1.64 standard deviations above its null expectation, and positive otherwise. Therefore, the null hypothesis is rejected if and only if this function is nonpositive.

We can code this function in R as follows.

```

682 # Rosenbaum (2018) objective: expectation - observed + kappa * SD
683 # Nonpositive values imply rejection at level alpha
684 sens_objective <- function(null_expect,
685                             null_variance,
686                             obs_stat,
687                             alpha = 0.05) {
688   # null_expect : overall null expectation of the test statistic
689   # null_variance : overall null variance of the test statistic
690   # obs_stat : observed value of the test statistic
691   # alpha : test size (e.g., 0.05 gives kappa about 1.64)
692
693   # Normal critical value
694   kappa <- qnorm(
695     p = 1 - alpha
696   )
697
698   # Null standard deviation
699   null_sd <- sqrt(
700     null_variance
701   )
702
703   # Gap between null expectation and observed statistic
704   # (null_expect - obs_stat), + Normal buffer (kappa * null_sd)
705   obj_val <- (null_expect - obs_stat) + kappa * null_sd
706
707   return(obj_val)
708 }

```

To build intuition, we next evaluate this function at $\Gamma = 1$.

```

709 # Evaluate under Gamma = 1 using null_ev, null_var, and obs_stat
710 sens_objective(
711   null_expect = null_ev,
712   null_variance = null_var,
713   obs_stat = obs_stat,
714   alpha = 0.05
715 )

```

[1] -0.06500212

The resulting value is negative, indicating rejection of the null at the chosen significance level. This is exactly what we would expect: Because we previously rejected the sharp null hypothesis that $\tau_h = 0$ for all units at $\Gamma = 1$, the function evaluated at $\Gamma = 1$ is indeed nonpositive.

The crucial property of this function is that it is concave over $\mathcal{U}_+$, which means that a tangent line drawn at *any* $\mathbf{u} \in \mathcal{U}_+$ lies above the function *everywhere* on $\mathcal{U}_+$. Rosenbaum (2018) draws the tangent line at the configuration of $\mathbf{u}$ chosen by the separable approximation. Because the original function is concave, this tangent line provides a global upper bound on the function for all candidate configurations of hidden confounding. If this upper bound is nonpositive, then the function itself must also be nonpositive. Consequently, rather than maximizing a curved function over $\mathcal{U}_+$, we can instead maximize a linear function, which greatly simplifies the optimization.

Straight lines are easy to work with because they decompose additively across matched sets. We can therefore determine, within each stratum separately, which hidden-bias pattern makes the tangent line as large as possible, and then add up these stratum-specific contributions. The resulting sum is the maximum value of the tangent line over all $\mathbf{u} \in \mathcal{U}_+$.

Recalling that the original function is concave, if the maximum value of the tangent-line approximation over $\mathbf{u} \in \mathcal{U}_+$ is at or below zero, then the original function itself must be nonpositive for every $\mathbf{u} \in \mathcal{U}_+$. Therefore, we can safely base the decision to reject the null hypothesis or not on the Taylor approximation. Although valid, this procedure can be conservative: The tangent line may remain above zero even when the original function would be below zero at its true worst-case $\mathbf{u}$. By contrast, the separable approximation can be liberal, as it may reject even when the worst-case value of the function is positive.

Conducting Worst-Case Sensitivity Analysis Below we report the upper one-sided p -value obtained under the separable approximation, along with an indication of whether the resulting decision to reject or not reject the null hypothesis agrees with the decision implied by the Taylor series approximation. Agreement or disagreement between the two methods is determined by the sign of the Taylor-series-based bound on the function described above that governs rejection of the null hypothesis. Disagreement can occur only when the separable approximation yields a p -value slightly below α , indicating rejection, while the Taylor series approximation yields a positive value, leading to non-rejection.

```

716 # Grid of Gamma values
717 Gamma_vals <- seq(
718   from = 1,
719   to = 1.5,
720   by = 0.0001
721 )
722
723 # lapply() runs sensitivity analysis at each Gamma, returns list of results
724 sens_results <- lapply(
725   X = Gamma_vals,
726   FUN = function(g) {
727     out <- senstrat(

```

```

728     sc      = data_matched$ldur_tilde_hm_scaled, # outcome
729     z       = data_matched$UN,                  # treatment indicator
730     st      = data_matched$fm,                  # matched set
731     gamma   = g,                               # Gamma
732     alternative = "greater",                    # upper-tail test
733     detail  = TRUE # return intermediate quantities
734   )
735
736   # Separable p-value
737   p_sep <- as.numeric(
738     out$Separable["P-value"]
739   )
740
741   # Taylor series margin (cvA)
742   cvA <- as.numeric(
743     out$lambda["Linear Taylor bound"]
744   )
745
746   # Decisions
747   sep_reject <- p_sep <= alpha                # separable rejects?
748   tay_reject <- cvA <= 0                      # Taylor rejects?
749
750   # Store results for this value of Gamma
751   data.frame(
752     Gamma      = g,          # sensitivity parameter
753     p_sep      = p_sep,      # separable p-value
754     cvA        = cvA,        # Taylor bound
755     agree_tay  = (sep_reject == tay_reject) # agreement indicator
756   )
757 }
758 )
759
760 # Bind into a single data frame
761 sens_df <- do.call(
762   what = rbind,
763   args = sens_results
764 )
765
766 # Smallest Gamma where the separable p-value is >= alpha
767 sens_value_sep <- min(
768   sens_df$Gamma[sens_df$p_sep >= alpha]
769 )
770
771 # Taylor series sensitivity value:
772 # smallest Gamma at which Taylor margin becomes positive (non-rejection)
773 sens_value_tay <- min(
774   sens_df$Gamma[sens_df$cvA > 0]
775 )

```

For this matched dataset, the conclusions from the separable and Taylor series approximations agree for nearly all values of Γ , differing over only a narrow interval from 1.1361 to 1.1399. As these results indicate, we can reject the sharp null of no effect under as-if randomization (i.e., $\Gamma = 1$), but this conclusion is sensitive to departures from that assumption. Using the more conservative Taylor series approximation, we can no longer reject the sharp null

of no effect against the alternative of a larger effect at the $\alpha = 0.05$ level once Γ reaches approximately 1.14.

The plot below illustrates this comparison and clarifies how the two approaches relate across values of Γ .




Figure 3: Sensitivity analysis for the sharp null of no individual treatment effect for all units. The top panel shows upper one-sided p -values under the separable approximation across values of Γ ; the bottom panel reports, across the same values of Γ , the corresponding Taylor series margin, with the zero line marking the rejection boundary. The shaded vertical band highlights values of Γ for which the separable approximation and the Taylor series approximation lead to different rejection decisions.

Zooming in on the shaded vertical band — corresponding to values of Γ for which the two approaches yield different conclusions (between 1.1361 to 1.1399) — shows that the two methods diverge in exactly the way one would expect. Over this range of Γ values, the Taylor series margin lies just above zero, indicating no rejection of the sharp null, while the separable approximation yields p -values just below 0.05, indicating rejection. This pattern reflects the relative conservatism of the Taylor series approximation: Rejection under the Taylor series approach implies rejection under the separable approach, but not vice versa.

Both approaches to sensitivity analysis rely on multiple approximations. For example, the separable approximation is used to approximate the worst-case configuration of hidden

bias $\mathbf{u} \in \mathcal{U}_+$, and inference additionally relies on a Normal approximation to the null randomization distribution of the test statistic. If the exact randomization distribution were used instead, the resulting sensitivity conclusions could differ for certain values of Γ . In principle, one could also assess sensitivity at a given Γ using the exact randomization distribution by fixing a configuration of hidden bias $\mathbf{u}$ — for example, the separable choice — and deriving, from (3), the implied probability distribution on Ω , which can then be combined with the previously defined `exact_sharp_null_dist` to compute the corresponding p -value.

2.4.2 HOW DO INFERENCES UNDER THE WEAK FRAMEWORK CHANGE UNDER THESE DEPARTURES?

Compared to tests of the sharp null hypothesis of no effect, constructing valid tests of an analogous weak null hypothesis is more challenging. A sharp null hypothesis specifies all missing potential outcomes, so we need only consider p -values over configurations of $\mathbf{u} \in \mathcal{U}_+$, rather than also over multiple configurations of potential outcomes. By contrast, a weak null permits many possible configurations of potential outcomes, so identifying the worst-case p -value requires optimization over both $\mathbf{u}$ and the potential outcomes consistent with the null. This joint optimization is computationally tractable only in restricted settings, such as those with binary outcomes (Fogarty et al., 2017) or matched-pair designs (Fogarty, 2020).

A more generally applicable approach is proposed by Fogarty (2023), which avoids the need for an explicit search for the worst-case p -value over all configurations of $\mathbf{u}$ and potential outcomes consistent with the null hypothesis about τ . To build intuition for this approach, note that the following inverse probability weighted (IPW) version of the Difference-in-Means is unbiased for the ATE, τ :

$$\sum_{s=1}^S (n_s/n) \left(\frac{1}{|\Omega_s|} \right) \left(\frac{\hat{\tau}_s}{p(\mathbf{z}_s)} \right), \quad (5)$$

where $p(\mathbf{z}_s)$ denotes the probability of assignment $\mathbf{z}_s$ among all treatment assignments in set s that hold fixed the realized number of treated units, with $|\Omega_s|$ denoting, as defined earlier, the total number of such assignments. Both $\hat{\tau}_s$ and $p(\mathbf{z}_s)$ are evaluated at the same realized treatment assignment $\mathbf{z}_s$ within matched set s . In other words, $p(\mathbf{z}_s)$ denotes the probability of the assignment under which the treated and control outcomes in set s — and hence the Difference-in-Means $\hat{\tau}_s$ — are realized.

When $\Gamma = 1$, this probability $p(\mathbf{z}_s)$ is $1/|\Omega_s|$, and the IPW Difference-in-Means reduces to the usual unweighted Difference-in-Means, $\hat{\tau}$. When $\Gamma > 1$, however, $p(\mathbf{z}_s)$ is no longer known exactly. In matched designs that may include sets with a single treated unit and many controls as well as sets with a single control unit and many treated units, Γ restricts the assignment probability to lie within bounds consistent with the specified degree of departure from as-if randomization:

$$\frac{1}{\Gamma(n_s - 1) + 1} \leq p(\mathbf{z}_s) \leq \frac{\Gamma}{(n_s - 1) + \Gamma}, \quad (6)$$

for all $\mathbf{z}_s \in \Omega_s$ and for all matched sets s .

Since the IPW Difference-in-Means in (5) cannot be directly computed when $\Gamma > 1$, Fogarty (2023) instead uses the bounds in (6) to construct a worst-case version of this IPW Difference-in-Means. For tests of a null hypothesis about the overall ATE across all matched sets, this worst-case procedure operates on the set-specific Difference-in-Means centered at the null ATE, $\hat{\tau}_s - \tau_0$, and applies the upper bound from (6) whenever $\hat{\tau}_s \geq \tau_0$ and the lower bound whenever $\hat{\tau}_s < \tau_0$. When testing the null against a smaller alternative, the procedure reverses these substitutions.

To see the value of this procedure, consider testing a null hypothesis about the ATE against the alternative of a larger ATE. Fogarty (2023) shows that, for any $\Gamma \geq 1$, the expected value of this worst-case IPW Difference-in-Means — defined under whatever the true distribution on Ω consistent with that Γ happens to be — is always less than or equal to the null when it is true. Conversely, when testing against the alternative of a smaller ATE, this expectation is greater than or equal to the null when it is true. Importantly, these properties hold for all possible values of $\mathbf{u}$ and all configurations of potential outcomes that are consistent with the null hypothesis.

How does this property of the expected value ensures a valid test in sufficiently large studies? Sticking to the case of testing against a larger ATE, the intuition is straightforward: If the expected value of the worst-case IPW Difference-in-Means is less than or equal to the null when it is true, then the worst-case IPW Difference-in-Means tends to fall at or below the null rather than above it. In other words, the procedure intentionally “tilts” the worst-case IPW Difference-in-Means in the direction opposite the alternative, making the procedure conservative.

In addition to this property of the expected value, a second key ingredient for valid hypothesis testing concerns variance estimation. For any fixed $\Gamma \geq 1$, an analogue of the variance estimator from Fogarty (2018) (discussed above) consistently overestimates (or bounds above) the true variance of the worst-case IPW Difference-in-Means. As a result, using this variance estimator further reduces the magnitude of the standardized test statistic — the worst-case IPW Difference-in-Means, which is already centered at the null ATE, divided by the square root of the estimated variance — relative to what would be obtained using the true (but unknown) variance.

Together, these two properties work in the same direction:

- The expected value is at or below the null, keeping the center of the distribution at or shifted away from the upper tail, and
- the variance estimate tends to be too large, which pulls the standardized statistic closer to zero.

As a result, when the null is true, the probability that the standardized statistic falls in the upper tail of the standard Normal distribution — that is, the probability of falsely rejecting the null — remains at or below the nominal significance level. Hence, the test is conservative (valid) in large studies. Analogous reasoning applies when testing against the alternative of a smaller effect: When the null hypothesis is true, the probability that the standardized

value falls in the lower tail of the standard Normal distribution also stays at or below the nominal significance level.

To implement this approach from Fogarty (2023), we first source from the GitHub repo an R function that, for a fixed $\Gamma \geq 1$, computes the worst-case IPW Difference-in-Means for a single matched set.

```

776 # Load the worst_case_IPW() function from GitHub repo
777 source(
778   paste0(
779     base_url, "/R/worst_case_IPW.R"
780   )
781 )

```

We now use this function to calculate the worst-case IPW version of the Difference-in-Means within each matched set. To form the overall statistic, we average the set-specific values using weights equal to each set's contribution to the total number of units. Finally, we center the worst-case IPW Difference-in-Means by the null ATE, standardize this difference using the conservative variance estimator from Fogarty (2018), and compare the resulting standardized test statistic to the standard Normal distribution to obtain p -values.

```

782 # One-sided sensitivity analysis for the weak null (ATE = 0)
783
784 # Null and alternative for the weak-null test
785 null_ATE <- 0 # Null: ATE = 0
786
787 alternative <- "greater" # Alt: ATE > 0 (upper-tail test)
788
789 # Data frame to store results for each Gamma
790 weak_sens_df <- data.frame(
791   Gamma = Gamma_vals,
792   p_value = NA_real_
793 )
794
795 for (g in Gamma_vals) { # Loop over Gamma values
796
797   # Worst-case IPW estimate within each matched set at Gamma = g
798   set_stats <- data_matched |>
799     group_by(
800       fm
801     ) |>
802     summarise(
803       n = n(), # Set size n_s
804       weight = n / nrow(data_matched), # Weight n_s / n
805       est = worst_case_IPW( # Worst-case IPW in set s
806         z = UN, # Treatment indicator
807         y = ldur, # Outcome
808         Gamma = g, # Sensitivity parameter (Gamma)
809         tau_h = null_ATE, # Weak null: ATE = 0
810         alternative = alternative # Direction of alternative
811       ),
812       .groups = "drop"
813     )
814

```



```

815 # Overall worst-case IPW statistic (weighted average across sets)
816 wc_ipw_stat <- sum(
817   set_stats$weight * set_stats$est
818 )
819
820 # Variance estimate for the worst-case IPW statistic
821 var_hat <- fine_strat_var_est(
822   strat_ns    = set_stats$n,
823   strat_ests = set_stats$est
824 )
825
826 # Standardized worst-case IPW statistic, centered at the null
827 se_hat <- sqrt(
828   var_hat
829 )
830
831 std_stat <- wc_ipw_stat / se_hat
832
833 # One-sided p-value, determined by the alternative
834 if (alternative == "greater") {
835   # Upper-tail test: ATE > 0
836   p_val <- 1 - pnorm(std_stat)
837 } else if (alternative == "less") {
838   # Lower-tail test: ATE < 0
839   p_val <- pnorm(std_stat)
840 } else {
841   stop("Only one-sided alternatives (`greater' or `less').")
842 }
843
844 # Store result for this Gamma
845 weak_sens_df$p_value[weak_sens_df$Gamma == g] <- p_val
846 }
847
848 # Sensitivity value: smallest Gamma where one-sided test no longer rejects
849 weak_sens_value <- min(
850   weak_sens_df$Gamma[weak_sens_df$p_value >= alpha]
851 )

```

Below we plot the upper one-sided p -values corresponding to Γ values from 1 to 1.5.




Figure 4: One-sided upper p -values for the weak null hypothesis of zero average effect across values of the sensitivity parameter Γ .

Under as-if randomization ($\Gamma = 1$), we find evidence of a positive average effect. This conclusion about the weak null is slightly more robust to departures from as-if randomization than our earlier rejection of the sharp null of no effect at $\Gamma = 1$. However, in absolute terms, the evidence for a positive ATE remains sensitive. Our qualitative conclusion about the null hypothesis that the ATE equals zero changes once Γ reaches 1.1756.

It is important to emphasize that this sensitivity analysis guarantees validity across *all* configurations of potential outcomes consistent with the weak null. By contrast, Fogarty (2023) also describes an alternative sensitivity analysis that restricts attention to subsets of configurations that researchers may regard as more plausible. Such alternatives can produce smaller p -values, but at the cost of forfeiting validity for certain (possibly unrealistic) potential outcome configurations.

3. Conclusion

We have now completed the full matching pipeline. We began by constructing matched sets and then conducted inference under the as-if randomization assumption, treating the matched design as a collection of completely randomized experiments within blocks. We examined inference under both sharp and weak null frameworks — corresponding, respectively, to unit-level causal effects and average effects — and then assessed the sensitivity of these

inferences to potential violations of the as-if randomization assumption. The embedded code, accompanied by extensive comments and illustrated through the running example from Gilligan and Sergenti (2008), is intended to be readily adapted by practitioners to their own datasets.

3.1 Limitations and Related Topics Beyond Our Scope

For clarity and focus, we have excluded several important topics, though we have pointed to relevant references where appropriate. Some of these omissions are not specific to matching or observational studies. In particular, we have not addressed imperfect compliance with treatment assignment, missing outcomes, clustered (rather than individual-level) assignment, or interference settings in which a subject’s outcome depends on others’ treatment assignments.

Our analysis has focused on inference for two causal targets only: a constant, additive treatment effect and the average treatment effect. We have not considered alternative effect models, such as dilated effects (Rosenbaum, 1999), multiplicative effects, or Tobit effects (Rosenbaum, 2010, pp. 46–49). Throughout, we have treated causal targets as fixed quantities (over the set of treatment assignments consistent with the matched design), rather than as random variables. As a result, we have excluded attributable effects (Rosenbaum, 2001, 2003; Hansen and Bowers, 2009) and the average treatment effect on the treated (Sekhon and Shem-Tov, 2021). We have also set aside methods for joint inference on sharp and weak causal hypotheses (Ding, 2017; Wu and Ding, 2021; Cohen and Fogarty, 2022), as well as strategies that improve the power of hypothesis tests by rescaling outcomes or choosing alternative test statistics, including regression-assisted approaches (Lin, 2013; Cohen and Fogarty, 2023; Guo and Basse, 2023).

Several omissions are specific to matching methodology itself:

- First, we have not covered a range of **alternative matching and balancing approaches**, including nonbipartite matching for multivalued treatments (Lu et al., 2001, 2011; Rosenbaum, 2010, pp. 207–221), template matching (Silber et al., 2014), multilevel matching (Zubizarreta and Keele, 2017; Pimentel et al., 2018), risk-set matching (Li et al., 2001), cardinality matching (Zubizarreta et al., 2014), coarsened exact matching (Iacus et al., 2012), exterior matching (Rosenbaum and Silber, 2013), or generalized full matching (Sävje et al., 2021), among others. We have further assumed that researchers specify *ex ante* the set of covariates to be balanced, thereby excluding screening procedures that identify additional covariates for adjustment based on observed imbalances (Cochran, 1965), as in tapered matching approaches (Daniel et al., 2008; Yu et al., 2021). Finally, we have not considered balance constraints such as fine and near-fine balance (Rosenbaum et al., 2007; Yang et al., 2012; Yu, 2023; Pimentel et al., 2015a,b), which enforce exact or nearly exact equality of covariate distributions.
- Second, we have not discussed the use of **prognostic covariates** identified using pilot or external data. Such approaches can substantially improve efficiency by prioritizing covariates that strongly predict potential outcomes. For theoretical foundations, see

Hansen (2008a), with complementary theoretical results and practical insights in Sales et al. (2018).

- Third, we have considered only matching strategies aimed at justifying as-if randomization, not those designed explicitly to improve **design sensitivity** — that is, robustness to moderate violations of as-if randomization (see Rosenbaum, 2010, Part III, pp. 257-311). Beyond match construction, researchers can also enhance design sensitivity by selecting particular test statistics, among other strategies (Rosenbaum, 2004; Heller et al., 2009; Hsu et al., 2013; Small et al., 2013).
- Fourth, we have excluded **extensions to covariate balance testing** beyond the framework developed in Hansen and Bowers (2008), including subsequent contributions by Gagnon-Bartsch and Shem-Tov (2019) and Branson (2021), as well as related approaches based on the stepwise intersection–union principle (Hansen and Sales, 2015).
- Fifth, we have ignored **residual imbalance on observed covariates**. Specifically, we have proceeded under the assumption that matched set membership can justify as-if randomization, even when small imbalances remain within matched sets. Under this approach, concerns about residual imbalance are absorbed into the sensitivity analysis for hidden bias. This approach is not entirely satisfying, however, because imbalances on observed covariates are not, in fact, hidden. Several methods therefore address such imbalances directly, either in approaches targeting individual effects (Rosenbaum, 1988; Pimentel and Huang, 2024; Chen et al., 2023; Heng et al., 2025) or average effects (Zhu et al., 2025).
- Sixth, we have not incorporated into the overall pipeline an explicit approach for defining an **interpretable study population** in matched designs. Common practice excludes units based on limited overlap in the estimated propensity score (Crump et al., 2009; King and Zeng, 2006), but study populations defined in this way can be difficult to characterize substantively. Hence, alternative approaches define the study population directly in terms of a small number of covariates. For example, Traskin and Small (2011) use a classification tree to define a study population that approximates one defined by trimming on the estimated propensity score, while Fogarty et al. (2016) use discrete optimization to define a maximal, rectangular region of covariate space with adequate overlap on key covariates.
- Seventh, we have not considered **quasi-experimental devices** (Campbell, 1957; Campbell and Stanley, 1963; Shadish et al., 2002), specifically how matched designs can leverage such devices (see Rosenbaum, 2025, Section IV). Two prominent quasi-experimental devices include an additional control group (Rosenbaum, 1987b) and outcomes with known effects (Rosenbaum, 1989), each of which has a range of valuable applications. In the context of evidence factors (Rosenbaum, 2021), one use of a second control group is to yield stronger causal evidence that is less susceptible to hidden confounding (Rosenbaum, 2023b). For known effects, one illustrative use is to support tests for hidden confounding that inform sensitivity analyses by ruling out certain violations of as-if randomization (Rosenbaum, 2023a).

- Eighth and finally, we have not addressed settings in which matched set membership itself may depend on the treatment assignments — that is, settings with **Z-dependence** (Pashley et al., 2021; Pimentel and Huang, 2024; Pimentel and Yu, 2024). Our exposition has relied on an analogy to Rubin’s framework of “assignment to treatment on the basis of a covariate” (Rubin, 1977), where matched set membership plays the role of the covariate. This analogy treats set membership as fixed, but in practice membership can be random when the realized treatment assignment — determined prior to matching — induces the matched structure, implying that the structure may vary across possible assignments. We have set aside this issue, which may be minor when practitioners use sufficiently tight calipers for matching (a point we thank Professor Samuel Pimentel for emphasizing in helpful discussions).

Despite these limitations, the pipeline and design-based perspective developed here provide a coherent framework for understanding how matched designs support causal inference. The emphasis on design-based foundations clarifies the logic underlying the construction of matched sets and the subsequent steps of inference and sensitivity analysis under both weak and sharp frameworks. We present this work not as exhaustive, but as a foundation on which researchers can build when incorporating more specialized methods tailored to their substantive and inferential goals.

References

- A. Albert and J.A. Anderson. On the existence of maximum likelihood estimates in logistic regression models. *Biometrika*, 71(1):1–10, 1984.
- Peter M. Aronow, Donald P. Green, and Donald K.K. Lee. Sharp bounds on the variance in randomized experiments. *The Annals of Statistics*, 42(3):850–871, 2014.
- Peter C. Austin. A critical appraisal of propensity-score matching in the medical literature between 1996 and 2003. *Statistics in Medicine*, 27(12):2037–2049, 2008.
- Peter C. Austin. Balance diagnostics for comparing the distribution of baseline covariates between treatment groups in propensity-score matched samples. *Statistics in Medicine*, 28(25):3083–3107, 2009.
- Zach Branson. Randomization tests to assess covariate balance when designing and analyzing matched datasets. *Observational Studies*, 7(2):1–36, 2021.
- Donald T. Campbell. Factors relevant to the validity of experiments in social settings. *Psychological Bulletin*, 54(4):297–312, 1957.
- Donald T. Campbell and Julian C. Stanley. *Experimental and Quasi-Experimental Designs for Research*. Houghton Mifflin Company, Boston, MA, 1963.
- Devin Caughey, Allan Dafoe, Xinran Li, and Luke Miratrix. Randomisation inference beyond the sharp null: Bounded null hypotheses and quantiles of individual treatment effects. *Journal of the Royal Statistical Society Series B (Statistical Methodology)*, 85(5):1471–1491, 2023.

- Kan Chen, Siyu Heng, Qi Long, and Bo Zhang. Testing biased randomization assumptions and quantifying imperfect matching and residual confounding in matched observational studies. *Journal of Computational and Graphical Statistics*, 32(2):528–538, 2023.
- Eun Yi Chung and Joseph P. Romano. Exact and asymptotically robust permutation tests. *The Annals of Statistics*, 41(2):484–507, 2013.
- William G. Cochran. The planning of observational studies of human populations. *Journal of the Royal Statistical Society. Series A (General)*, 128(2):234–266, 1965.
- William G. Cochran and Donald B. Rubin. Controlling bias in observational studies: A review. *Sankhyā: The Indian Journal of Statistics, Series A (1961–2002)*, 35(4):417–446, 1973.
- Peter L. Cohen and Colin B. Fogarty. Gaussian pre pivoting for finite population causal inference. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 84(2):295–320, 2022.
- Peter L. Cohen and Colin B. Fogarty. No-harm calibration for generalized Oaxaca-Blinder estimators. *Biometrika*, 111(1):331–338, 2023.
- David R. Cox. *Planning of Experiments*. Wiley, New York, NY, 1958.
- Richard K. Crump, V. Joseph Hotz, Guido W. Imbens, and Oscar A. Mitnik. Dealing with limited overlap in estimation of average treatment effects. *Biometrika*, 96(1):187–199, 2009.
- Shoshana R. Daniel, Katrina Armstrong, Jeffrey H. Silber, and Paul R. Rosenbaum. An algorithm for optimal tapered matching, with application to disparities in survival. *Journal of Computational and Graphical Statistics*, 17(4):914–924, 2008.
- Peng Ding. A paradox from randomization-based causal inference. *Statistical Science*, 32(3):331–345, 2017.
- Ronald A. Fisher. *The Design of Experiments*. Oliver and Boyd, Edinburgh, SCT, 1935.
- Colin B. Fogarty. On mitigating the analytical limitations of finely stratified experiments. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 80(5):1035–1056, 2018.
- Colin B. Fogarty. Studentized sensitivity analysis for the sample average treatment effect in paired observational studies. *Journal of the American Statistical Association*, 115(531):1518–1530, 2020.
- Colin B. Fogarty. Testing weak nulls in matched observational studies. *Biometrics*, 79(3):2196–2207, 2023.
- Colin B. Fogarty, Mark E. Mikkelsen, David F. Gaieski, and Dylan S. Small. Discrete optimization for interpretable study populations and randomization inference in an observational study of severe sepsis mortality. *Journal of the American Statistical Association*, 111(514):447–458, 2016.

- Colin B. Fogarty, Pixu Shi, Mark E. Mikkelsen, and Dylan S. Small. Randomization inference and sensitivity analysis for composite null hypotheses with binary outcomes in matched observational studies. *Journal of the American Statistical Association*, 112(517):321–331, 2017.
- Johann Gagnon-Bartsch and Yotam Shem-Tov. The classification permutation test: A flexible approach to testing for covariate imbalance in observational studies. *The Annals of Applied Statistics*, 13(3):1464–1483, 2019.
- Joseph L. Gastwirth, Abba M. Krieger, and Paul R. Rosenbaum. Asymptotic separability in sensitivity analysis. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 63(3):545–555, 2000.
- Andrew Gelman. A Bayesian formulation of exploratory data analysis and goodness-of-fit testing. *International Statistical Review*, 71(2):369–382, 2003.
- Andrew Gelman. Review: Causality and statistical learning. *American Journal of Sociology*, 117(3):955–966, 2011.
- Michael J. Gilligan and Ernest J. Sergenti. Do UN interventions cause peace? Using matching to improve causal inference. *Quarterly Journal of Political Science*, 3(2):89–122, 2008.
- Xing Sam Gu and Paul R. Rosenbaum. Comparison of multivariate matching methods: Structures, distances, and algorithms. *Journal of Computational and Graphical Statistics*, 2(4):405–420, 1993.
- Kevin Guo and Guillaume W. Basse. The generalized Oaxaca-Blinder estimator. *Journal of the American Statistical Association*, 118(541):524–536, 2023.
- Ben B. Hansen. Full matching in an observational study of coaching for the SAT. *Journal of the American Statistical Association*, 99(467):609–618, 2004.
- Ben B. Hansen. The prognostic analogue of the propensity score. *Biometrika*, 95(2):481–488, 2008a.
- Ben B. Hansen. The essential role of balance tests in propensity-matched observational studies: Comments on ‘a critical appraisal of propensity-score matching in the medical literature between 1996 and 2003’ by Peter Austin. *Statistics in Medicine*, 27(12):2050–2054, 2008b.
- Ben B. Hansen. Propensity score matching to extract latent experiments from nonexperimental data: A case study. In Neil J. Dorans and Sandip Sinharay, editors, *Looking Back: Proceedings of a Conference in Honor of Paul W. Holland*, volume 202 of *Lecture Notes in Statistics*, chapter 9, pages 149–181. Springer, New York, NY, 2011.
- Ben B. Hansen and Jake Bowers. Covariate balance in simple, stratified and clustered comparative studies. *Statistical Science*, 23(2):219–236, 2008.
- Ben B. Hansen and Jake Bowers. Attributing effects to a cluster-randomized get-out-the-vote campaign. *Journal of the American Statistical Association*, 104(487):873–885, 2009.

- Ben B. Hansen and Stephanie Olsen Klopfer. Optimal full matching and related designs via network flows. *Journal of Computational and Graphical Statistics*, 15(3):609–627, 2006.
- Ben B. Hansen and Adam Sales. Comment on cochran’s “observational studies”. *Observational Studies*, pages 184–193, 2015.
- Christopher Harshaw, Joel A. Middleton, and Fredrik Sävje. Optimized variance estimation under interference and complex experimental designs. arXiv Preprint, <https://arxiv.org/pdf/2112.01709>, July 2024.
- Ruth Heller, Paul R. Rosenbaum, and Dylan S. Small. Split samples and design sensitivity in observational studies. *Journal of the American Statistical Association*, 104(487):1090–1101, 2009.
- Siyu Heng, Yanxin Shen, and Pengyun Wang. Reconciling overt bias and hidden bias in sensitivity analysis for matched observational studies. arXiv Preprint, <https://arxiv.org/pdf/2311.11216>, May 2025.
- Joseph L. Hodges Jr. and Erich L. Lehmann. Estimates of location based on rank tests. *The Annals of Mathematical Statistics*, 34(2):598–611, 1963.
- Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.
- Jesse Y. Hsu, Dylan S. Small, and Paul R. Rosenbaum. Effect modification and design sensitivity in observational studies. *Journal of the American Statistical Association*, 108(501):135–148, 2013.
- Stefano M. Iacus, Gary King, and Giuseppe Porro. Causal inference without balance checking: Coarsened exact matching. *Political Analysis*, 20(1):1–24, 2012.
- Kosuke Imai, Gary King, and Elizabeth A. Stuart. Misunderstandings between experimentalists and observationalists about causal inference. *Journal of the Royal Statistical Society: Series A (Statistics in Society)*, 171(2):481–502, 2008.
- Gary King and Langche Zeng. The dangers of extreme counterfactuals. *Political Analysis*, 14(2):131–159, 2006.
- Fan Li, Kari Lock Morgan, and Alan M. Zaslavsky. Balancing covariates via propensity score weighting. *Journal of the American Statistical Association*, 113(521):390–400, 2018.
- Yunfei Paul Li, Kathleen J. Propert, and Paul R. Rosenbaum. Balanced risk set matching. *Journal of the American Statistical Association*, 96(455):870–882, 2001.
- Winston Lin. Agnostic notes on regression adjustments to experimental data: Reexamining Freedman’s critique. *The Annals of Applied Statistics*, 7(1):295–318, 2013.
- Bo Lu, Elaine Zanutto, Robert Hornik, and Paul R. Rosenbaum. Matching with doses in an observational study of a media campaign against drug abuse. *Journal of the American Statistical Association*, 96(456):1245–1253, 2001.

- Bo Lu, Robert Greevy, Xinyi Xu, and Cole Beck. Optimal nonbipartite matching and its statistical applications. *The American Statistician*, 65(1):21–30, 2011.
- Prasanta C. Mahalanobis. On the generalized distance in statistics. *Proceedings of Indian National Science Academy*, 2(1):49–55, 1936.
- Luke W. Miratrix, Jasjeet S. Sekhon, and Bin Yu. Adjusting treatment effect estimates by post-stratification in randomized experiments. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 75(2):369–396, 2013.
- Jersey Neyman. Sur les applications de la théorie des probabilités aux expériences agricoles: Essai des principes. *Roczniki Nauk Rolniczych*, 10:1–51, 1923.
- Nicole E. Pashley and Luke W. Miratrix. Insights on variance estimation for blocked and matched pairs designs. *Journal of Educational and Behavioral Statistics*, 46(3):271–296, 2021.
- Nicole E. Pashley, Guillaume W. Basse, and Luke W. Miratrix. Conditional as-if analyses in randomized experiments. *Journal of Causal Inference*, 9(1), 2021.
- Samuel D. Pimentel and Yaxuan Huang. Covariate-adaptive randomization inference in matched designs. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 86(5):1312–1338, 2024.
- Samuel D. Pimentel and Ruoqi Yu. Re-evaluating the impact of hormone replacement therapy on heart disease using match-adaptive randomization inference. Working Paper, <https://arxiv.org/pdf/2403.01330.pdf>, March 2024.
- Samuel D. Pimentel, Rachel R. Kelz, Jeffrey H. Silber, and Paul R. Rosenbaum. Large, sparse optimal matching with refined covariate balance in an observational study of the health outcomes produced by new surgeons. *Journal of the American Statistical Association*, 110(510):515–527, 2015a.
- Samuel D. Pimentel, Frank Yoon, and Luke Keele. Variable-ratio matching with fine balance in a study of the Peer Health Exchange. *Statistics in Medicine*, 34(30):4070–4082, 2015b.
- Samuel D. Pimentel, Lindsay C. Page, Matthew Lenard, and Luke Keele. Optimal multilevel matching using network flows: An application to a summer reading intervention. *Annals of Applied Statistics*, 12(3):1479–1505, 2018.
- James M. Robins. Confidence intervals for causal parameters. *Statistics in Medicine*, 7(7):773–785, 1988.
- Paul R. Rosenbaum. Conditional permutation tests and the propensity score in observational studies. *Journal of the American Statistical Association*, 79(387):565–574, 1984.
- Paul R. Rosenbaum. Sensitivity analysis for certain permutation inferences in matched observational studies. *Biometrika*, 74(1):13–26, 1987a.
- Paul R. Rosenbaum. The role of a second control group in an observational study. *Statistical Science*, 2(3):292–306, 1987b.

- Paul R. Rosenbaum. Permutation tests for matched pairs with adjustments for covariates. *Journal of the Royal Statistical Society Series C: Applied Statistics*, 37(3):401–411, 1988.
- Paul R. Rosenbaum. The role of known effects in observational studies. *Biometrics*, 45(2): 557–569, 1989.
- Paul R. Rosenbaum. A characterization of optimal designs for observational studies. *Journal of the Royal Statistical Society. Series B (Methodological)*, 53(3):597–610, 1991.
- Paul R. Rosenbaum. Hodges–Lehmann point estimates of treatment effect in observational studies. *Journal of the American Statistical Association*, 88(424):1250–1253, 1993.
- Paul R. Rosenbaum. Quantiles in nonrandom samples and observational studies. *Journal of the American Statistical Association*, 90(432):1424–1431, 1995.
- Paul R. Rosenbaum. Reduced sensitivity to hidden bias at upper quantiles in observational studies with dilated treatment effects. *Biometrics*, 55(2):560–564, 1999.
- Paul R. Rosenbaum. Effects attributable to treatment: Inference in experiments and observational studies with a discrete pivot. *Biometrika*, 88(1):219–231, 2001.
- Paul R. Rosenbaum. Exact confidence intervals for nonconstant effects by inverting the signed rank test. *The American Statistician*, 57(2):132–138, 2003.
- Paul R. Rosenbaum. Design sensitivity in observational studies. *Biometrika*, 91(1):153–164, 2004.
- Paul R. Rosenbaum. *Design of Observational Studies*. Springer, New York, NY, 1st edition, 2010.
- Paul R. Rosenbaum. Optimal matching of an optimally chosen subset in observational studies. *Journal of Computational and Graphical Statistics*, 21(1):57–71, 2012.
- Paul R. Rosenbaum. *Observation and Experiment: An Introduction to Causal Inference*. Harvard University Press, Cambridge, MA, 2017.
- Paul R. Rosenbaum. Sensitivity analysis for stratified comparisons in an observational study of the effect of smoking on homocysteine levels. *Annals of Applied Statistics*, 12(4): 2312–2334, 2018.
- Paul R. Rosenbaum. *Replication and Evidence Factors in Observational Studies*. Number 167 in Monographs on Statistics and Applied Probability. CRC Press, Boca Raton, FL, 2021.
- Paul R. Rosenbaum. Sensitivity analyses informed by tests for bias in observational studies. *Biometrics*, 79(1):475–487, 2023a.
- Paul R. Rosenbaum. A second evidence factor for a second control group. *Biometrics*, 79(4): 3968–3980, 2023b.

- Paul R. Rosenbaum. *An Introduction to the Theory of Observational Studies*. Springer Texts in Statistics. Springer, New York, NY, 2025.
- Paul R. Rosenbaum and Abba M. Krieger. Sensitivity of two-sample permutation inferences in observational studies. *Journal of the American Statistical Association*, 85(410):493–498, 1990.
- Paul R. Rosenbaum and Donald B. Rubin. The bias due to incomplete matching. *Biometrics*, 41(1):103–116, 1985.
- Paul R. Rosenbaum and Jeffrey H. Silber. Using the exterior match to compare two entwined matched control groups. *The American Statistician*, 63(2):67–75, 2013.
- Paul R. Rosenbaum, Richard N. Ross, and Jeffrey H. Silber. Minimum distance matched sampling with fine balance in an observational study of treatment for ovarian cancer. *Journal of the American Statistical Association*, 102(477):75–83, 2007.
- Donald B. Rubin. Assignment to treatment group on the basis of a covariate. *Journal of Educational Statistics*, 2(1):1–26, 1977.
- Donald B. Rubin. Comment on ‘Randomization analysis of experimental data in the Fisher randomization test’ by Basu, D. *Journal of the American Statistical Association*, 75(371):591–593, 1980.
- Donald B. Rubin. Which ifs have causal answers? (Comment on ‘Statistics and causal inference’ by Paul W. Holland). *Journal of the American Statistical Association*, 81(396):961–962, 1986.
- Donald B. Rubin. Using propensity scores to help design observational studies: Application to the tobacco litigation. *Health Services and Outcomes Research Methodology*, 2(3):169–188, 2001.
- Donald B. Rubin and Richard P. Waterman. Estimating the causal effects of marketing interventions using propensity score methodology. *Statistical Science*, 21(2):206–222, 2006.
- Adam C. Sales, Ben B. Hansen, and Brian Rowan. Rebar: Reinforcing a matching estimator with predictions from high-dimensional covariates. *Journal of Educational and Behavioral Statistics*, 43(1):3–31, 2018.
- Fredrik Sävje, Michael J. Higgins, and Jasjeet S. Sekhon. Generalized full matching. *Political Analysis*, 29(4):423–447, 2021.
- Jasjeet S. Sekhon and Yotam Shem-Tov. Inference on a new class of sample average treatment effects. *Journal of the American Statistical Association*, 116(534):798–804, 2021.
- William R. Shadish, Thomas D. Cook, and Donald T. Campbell. *Experimental and Quasi-Experimental Designs for Generalized Causal Inference*. Houghton Mifflin Company, Boston, MA, 2002.

- Jeffrey H. Silber, Paul R. Rosenbaum, Richard N. Ross, Justin M. Ludwig, Wei Wang, Bijan A. Niknam, Nabanita Mukherjee, Philip A. Saynisch, Orit Even-Shoshan, Rachel R. Kelz, and Lee A. Fleisher. Template matching for auditing hospital cost and quality. *Health Services Research*, 48(5):1446–1474, 2014.
- Dylan S. Small, Jing Cheng, M. Elizabeth Halloran, and Paul R. Rosenbaum. Case definition and design sensitivity. *Journal of the American Statistical Association*, 108(504):1457–1468, 2013.
- Elizabeth A. Stuart. Matching methods for causal inference: A review and a look forward. *Statistical Science*, 25(1):1–21, 2010.
- Mikhail Traskin and Dylan S. Small. Defining the study population for an observational study to ensure sufficient overlap: A tree approach. *Statistics in Biosciences*, 3:94–118, 2011.
- Jason Wu and Peng Ding. Randomization tests for weak null hypotheses. *Journal of the American Statistical Association*, 116(536):1898–1913, 2021.
- Dan Yang, Dylan S. Small, Jeffrey H. Silber, and Paul R. Rosenbaum. Optimal matching with minimal deviation from fine balance in a study of obesity and surgical outcomes. *Biometrics*, 68(2):628–636, 2012.
- Ruoqi Yu. How well can fine balance work for covariate balancing. *Biometrics*, 79(3): 2346–2356, 2023.
- Ruoqi Yu, Dylan S. Small, and Paul R. Rosenbaum. The information in covariate imbalance in studies of hormone replacement therapy. *The Annals of Applied Statistics*, 15(4): 2023–2042, 2021.
- Jianan Zhu, Jeffrey Zhang, Zijian Guo, and Siyu Heng. Randomization-based inference for average treatment effects in inexactly matched observational studies. arXiv Preprint, <https://arxiv.org/pdf/2311.11216>, March 2025.
- José R. Zubizarreta and Luke Keele. Optimal multilevel matching in clustered observational studies: A case study of the effectiveness of private schools under a large-scale voucher system. *Journal of the American Statistical Association*, 112(518):547–560, 2017.
- José R. Zubizarreta, Ricardo D. Paredes, and Paul R. Rosenbaum. Matching for balance, pairing for heterogeneity in an observational study of the effectiveness of for-profit and not-for-profit high schools in Chile. *The Annals of Applied Statistics*, 8(1):204–231, 2014.